

Section 5

Deep Generative model

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

1

Acknowledgements

- The materials majorly derived
http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture14.pdf
- OpenAI spinningup
https://spinningup.openai.com/en/latest/spinningup/rl_intro.html
- David Silver
https://www.davidsilver.uk/wp-content/uploads/2020/03/intro_RL.pdf

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

2



PSL Overview

- Introduction
- [Fully visible belief network](#)
- [Boltzmann machine/RBM/DBM](#)
- [Autoencoder](#)
- [GAN](#)

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

3



PSL Supervised vs Unsupervised

Supervised Learning

Data: (x, y)
 x is data, y is label

Goal: Learn a *function* to map $x \rightarrow y$

Examples: Classification, regression, object detection, semantic segmentation, image captioning, etc.

Unsupervised Learning

Training data is cheap

Data: x
 Just data, no labels!

Holy grail: Solve unsupervised learning
 \Rightarrow understand structure of visual world

Goal: Learn some underlying hidden *structure* of the data

Examples: Clustering, dimensionality reduction, feature learning, density estimation, etc.

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

4

- Examples:
 - Dimension reduction : PCA
 - Clustering: k-means
 - Density estimation
 - Feature learning
- General framework:
 - Find deterministic function $f: z=f(x)$, x is data, z is the latent



Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

5

Unsupervised learning vs. Generative model

- $z = f(x)$ vs. $x = g(z)$
- $P(z|x)$ vs. $P(x|z)$
- **Encoder** vs. **Decoder (Generator)**
 - $P(x, z)$ needed. (cf : $P(y|x)$ in supervised learning)
 - $P(z|x) = P(x, z) / P(x)$
 - $P(x|z) = P(x, z) / P(z) \rightarrow P(z)$ is given. (prior)

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

6

Why generative models?

- Realistic samples for artwork, super-resolution, colorization, etc.



- Generative models of time-series data can be used for simulation and planning (reinforcement learning applications!)
- Training generative models can also enable inference of latent representations that can be useful as general features

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

7

Taxonomy of Generative Models

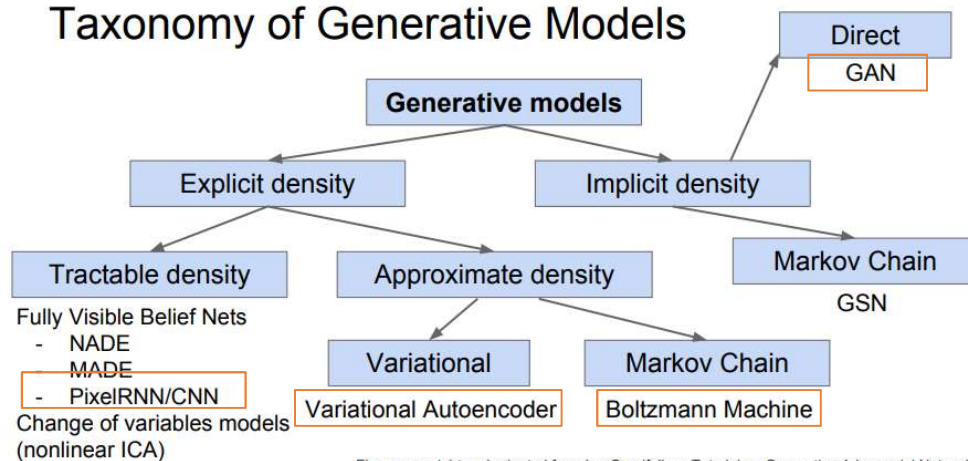


Figure copyright and adapted from Ian Goodfellow, Tutorial on Generative Adversarial Networks, 2017.

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

8



Fully visible belief network

Explicit density model

Use chain rule to decompose likelihood of an image x into product of 1-d distributions:

$$p(x) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

\uparrow Likelihood of image x \uparrow Probability of i 'th pixel value given all previous pixels

Will need to define ordering of "previous pixels"

Complex distribution over pixel values \Rightarrow Express using a neural network!

Then maximize likelihood of training data

Drawback: sequential generation is slow!


Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

9



PixelRNN [van der Oord et al. 2016]

- Generate image pixels starting from corner
- Dependency on previous pixels modeled using an RNN (LSTM)



Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

10

Still generate image pixels starting from corner

Dependency on previous pixels now modeled using a CNN over context region

Training is faster than PixelRNN
(can parallelize convolutions since context region values known from training images)

Generation must still proceed sequentially
=> still slow




Figure copyright van der Oord et al., 2016. Reproduced with permission.

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

11

Pros:

- Can explicitly compute likelihood $p(x)$
- Explicit likelihood of training data gives good evaluation metric
- Good samples

Con:

- Sequential generation => slow

Improving PixelCNN performance

- Gated convolutional layers
- Short-cut connections
- Discretized logistic loss
- Multi-scale
- Training tricks
- Etc...

See


- Van der Oord et al. NIPS 2016
- Salimans et al. 2017 (PixelCNN++)

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

12

What is Boltzmann machine?

- Boltzmann machines are stochastic and generative neural networks capable of learning internal representations
- They are able to represent (given sufficient time) and solve difficult combinatorial problems.
- They were invented in 1985 by Geoffrey Hinton
- non-deterministic (or stochastic) generative Deep Learning models with only two types of nodes: visible (v) and hidden (h)
- Unlike classical neural networks
 - No output nodes
 - Connection between input nodes (v)
- This allows them to share information among themselves and self-generate subsequent data




Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

13

Restricted Boltzmann machine (RBM)

- RBMs are a two-layered artificial neural network with generative capabilities.
- They have the ability to learn a probability distribution over its set of input.
- It can be used for dimensionality reduction, classification, regression, collaborative filtering, feature learning, and topic modeling.
- RBMs are a special class of Boltzmann Machines and they are restricted in terms of the connections between the visible and the hidden units.
- Every node in the visible layer is connected to every node in the hidden layer but no two nodes in the same group are connected to each other.
- This restriction allows for more efficient training algorithms than what is available for the general class of Boltzmann machines, in particular, the gradient-based contrastive divergence algorithm.





Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

14

$$h^1 = S(v^{(0)T}W + a)$$

- RBM is a Stochastic Neural Network which means that each neuron will have some random behavior when activated.
- There are two other layers of bias units (hidden bias and visible bias) in an RBM.
- The hidden bias RBM produce the activation on the forward pass and the visible bias helps RBM to reconstruct the input during a backward pass.
- The reconstructed input is always different from the actual input as there are no connections among the visible units and therefore, no way of transferring information among themselves.



$$v^1 = S(h^{(1)T}W^T + a)$$

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

15

- Consider the difference $v^{(0)} - v^{(1)}$ as reconstruction error. The weights are adjusted in each iteration so as to minimize this error
- In forward pass: $p(h^{(1)} | v^{(0)}; W)$
- In backward pass: $p(v^{(1)} | h^{(1)}; W)$
- It is a joint distribution: $p(v, h)$
- Assume that we have two normal distributions, one from the input data (denoted by $p(x)$) and one from the reconstructed input approximation (denoted by $q(x)$). The difference between these two distributions is our error in the graphical sense and our goal is to minimize it → Kullback-Leibler divergence (KL-divergence)
- KL-divergence measures the non-overlapping areas under the two graphs and the RBM's optimization algorithm tries to minimize this difference by changing the weights so that the reconstruction closely resembles the input.

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

16

Contrastive divergence

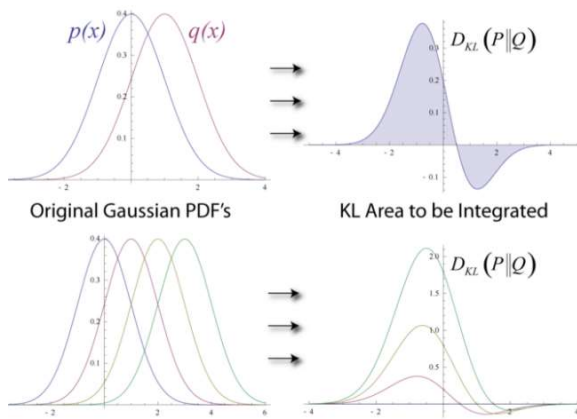


Image by [Mundhenk](#) on [Wikimedia](#)

Boltzmann Machines (and RBMs) are Energy-based models and a joint configuration, (\mathbf{v}, \mathbf{h}) of the visible and hidden units has an energy given by:

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i \in \text{visible}} a_i v_i - \sum_{j \in \text{hidden}} b_j h_j - \sum_{i,j} v_i h_j w_{ij}$$

The probability that the network assigns to a visible vector, \mathbf{v} , is given by summing over all possible hidden vectors

$$p(\mathbf{v}) = \frac{\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}} \quad \frac{\partial \log p(\mathbf{v})}{\partial w_{ij}} = \langle v_i h_j \rangle_{\text{data}} - \langle v_i h_j \rangle_{\text{model}}$$

$$\Delta w_{ij} = \alpha (\langle v_i h_j \rangle_{\text{data}} - \langle v_i h_j \rangle_{\text{model}})$$

Expectation
reconstruct

Learning rate

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines ParisTech, PSL

17

Contrastive divergence

- The learning rule is much more closely approximating the gradient of another objective function called the **Contrastive Divergence** which is the difference between two Kullback-Liebler divergences

Algorithm 1. k -step contrastive divergence

Input: RBM $(V_1, \dots, V_m, H_1, \dots, H_n)$, training batch S
Output: gradient approximation Δw_{ij} , Δb_j and Δc_i for $i = 1, \dots, n$, $j = 1, \dots, m$

```

1 init  $\Delta w_{ij} = \Delta b_j = \Delta c_i = 0$  for  $i = 1, \dots, n$ ,  $j = 1, \dots, m$ 
2 forall the  $\mathbf{v} \in S$  do
3    $\mathbf{v}^{(0)} \leftarrow \mathbf{v}$ 
4   for  $t = 0, \dots, k-1$  do
5     for  $i = 1, \dots, n$  do sample  $h_i^{(t)} \sim p(h_i | \mathbf{v}^{(t)})$ 
6     for  $j = 1, \dots, m$  do sample  $v_j^{(t+1)} \sim p(v_j | \mathbf{h}^{(t)})$ 
7   for  $i = 1, \dots, n$ ,  $j = 1, \dots, m$  do
8      $\Delta w_{ij} \leftarrow \Delta w_{ij} + p(H_i = 1 | \mathbf{v}^{(0)}) \cdot v_j^{(0)} - p(H_i = 1 | \mathbf{v}^{(k)}) \cdot v_j^{(k)}$ 
9      $\Delta b_j \leftarrow \Delta b_j + v_j^{(0)} - v_j^{(k)}$ 
10     $\Delta c_i \leftarrow \Delta c_i + p(H_i = 1 | \mathbf{v}^{(0)}) - p(H_i = 1 | \mathbf{v}^{(k)})$ 

```


Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines ParisTech, PSL

18

Deep belief network (DBM)

- Multiple RBMs can be stacked and can be fine-tuned through the process of gradient descent and back-propagation.
- One of first Deep-Learning models
- Proposed by G. Hinton in 2006
- Generative probabilistic model (mostly UNSUPERVISED)


For capturing high-order *correlations* of observed/visible data (\rightarrow pattern analysis, or synthesis); and/or characterizing *joint* statistical distributions of visible data



19

Use of trained RBM

- **Input data "completion" : set some v_i then compute h , and generate compatible full samples**
- **Generating representative samples**
- **Classification if trained with inputs=data+label**





Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

20

Interpretation of trained RBM hidden layer

- Look at weights of hidden nodes \rightarrow low-level features




This is why people are inspired to stack the RBMs to get more “abstract” features.

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines ParisTech, PSL

21

Learning of DBN

Greedy learning of successive layers



Algorithm 1 Recursive Greedy Learning Procedure for the DBN.


- 1: Fit parameters W^1 of the 1st layer RBM to data.
- 2: Freeze the parameter vector W^1 and use samples h^1 from $Q(h^1|v) = P(h^1|v, W^1)$ as the data for training the next layer of binary features with an RBM.
- 3: Freeze the parameters W^2 that define the 2nd layer of features and use the samples h^2 from $Q(h^2|h^1) = P(h^2|h^1, W^2)$ as the data for training the 3rd layer of binary features.
- 4: Proceed recursively for the next layers.

22

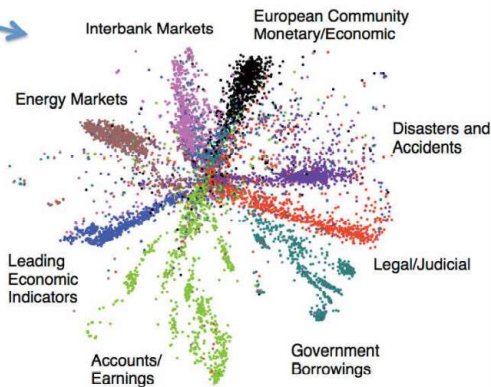


Example application of DBN: Clustering of documents in database

Model $P(\text{document})$



Reuters dataset: 804,414
newswire stories: **unsupervised**



*The goal of this model is to find the internal representation in all the types (classes) of documents provided by Reuters dataset

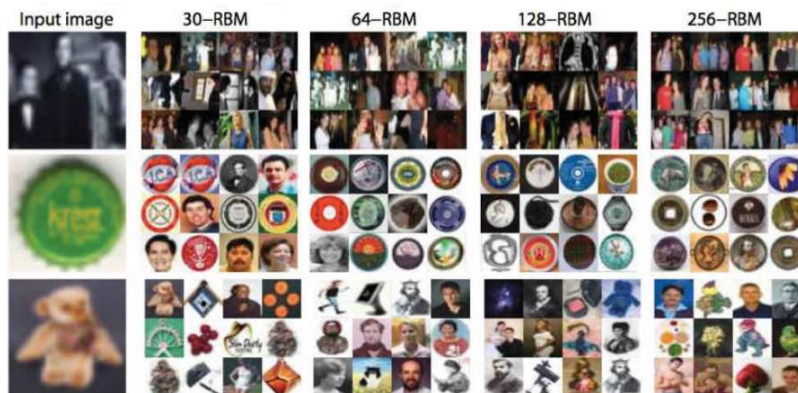
Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

23



Image Retrieval application example

- Map images in to binary codes fast retrieval




[Small codes, Torralba, Fergus, Weiss, CVPR 2008]

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

24

Autoencoders



Learn q_{θ} and p_{ϕ} in order to minimize reconstruction cost:

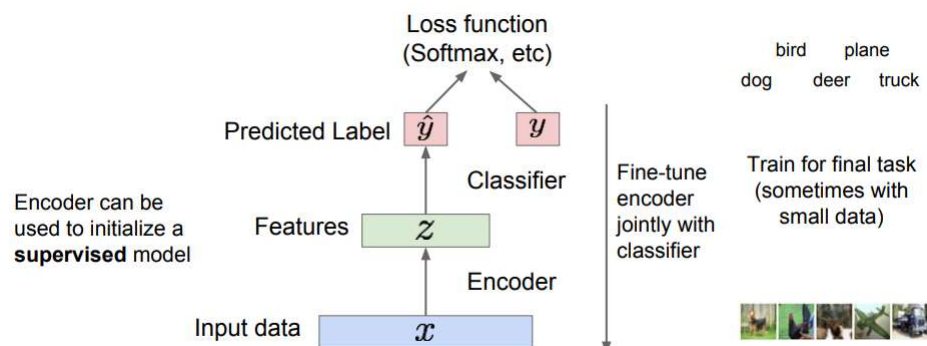

$$Q = \sum_k \|\hat{x}_k - x_k\|^2 = \sum_k \|p_{\phi}(q_{\theta}(x_k)) - x_k\|^2$$

→ unsupervised learning of latent variables,
and of a generative model

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines ParisTech, PSL

25


Use case of AutoEncoder



- • Denoising autoencoders
- Sparse autoencoders
- Stochastic autoencoders
- Contractive autoencoders
- VARIATIONAL autoencoders
- ...

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

27



Greedy layerwise training:

for each layer k , use backpropagation to minimize

$$\|A_k(h^{(k)}) - h^{(k)}\|^2 \quad (+ \text{regularization cost } \lambda \sum_{ij} |W_{ij}|^2)$$

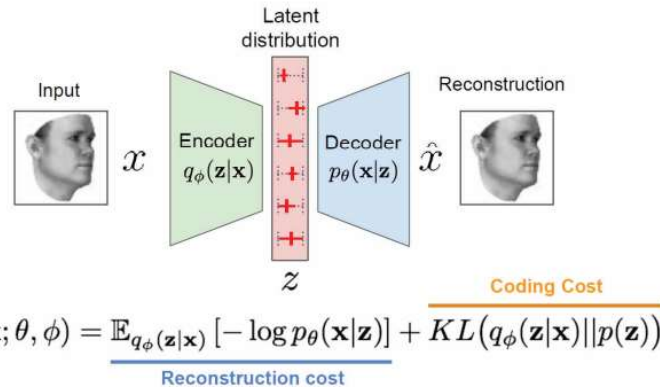
possibly + additional term for "sparsity"

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

28

Variational AutoEncoders (VAE)

Kingma et al. 2014
Rezende et al. 2014



Slide: Irina Higgins, Loïc Matthey

KL = Kullback-Leibler divergence (a.k.a. 'relative entropy')
KL(Q || P) measures how different are distributions

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

29

Generative Adversarial Network (GAN)

*[Introduced in 2014 by Ian Goodfellow et al.
(incl. Yoshua Bengio) from University of Montreal]*

Goal: generate « artificial » but credible examples
credible = sampled from same probability distribution $p(x)$

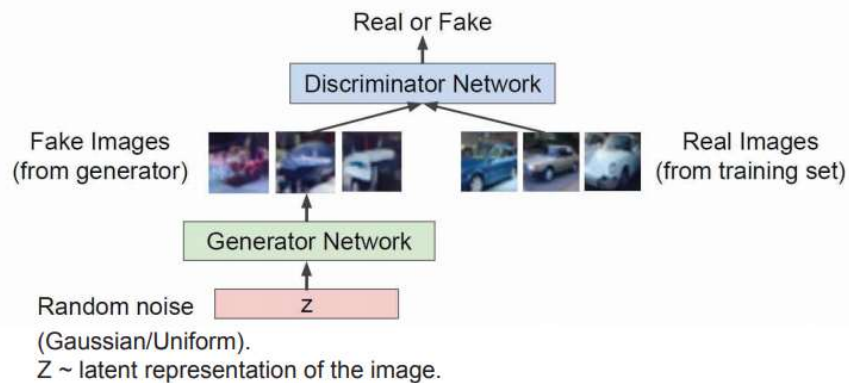
Idea: instead of trying to explicitly estimate $p(x)$,
 1. **LEARN** a transformation G from a simple and known distribution (e.g. random) into X ,
 2. then sampling $z \rightarrow$ generate realistic samples $G(z)$

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

30

Generator network: try to fool the discriminator by generating real-looking images

Discriminator network: try to distinguish between real and fake images



Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

31

$$\min_G \max_D V(D, G)$$

It is formulated as a **minimax game**, where:

- The Discriminator is trying to maximize its reward $V(D, G)$
- The Generator is trying to minimize Discriminator's reward (or maximize its loss)

$$V(D, G) = \mathbb{E}_{x \sim p(x)} [\log D(x)] + \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

Joint training of D and G

The Nash equilibrium of this particular game is achieved at:

- $P_{data}(x) = P_{gen}(x) \quad \forall x$
- $D(x) = \frac{1}{2} \quad \forall x$

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

32

In practice, alternate Discriminator training (gradient ascent) and Generator training:

for number of training iterations do

for k steps do

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m [\log D_{\theta_d}(x^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(z^{(i)})))]$$

end for

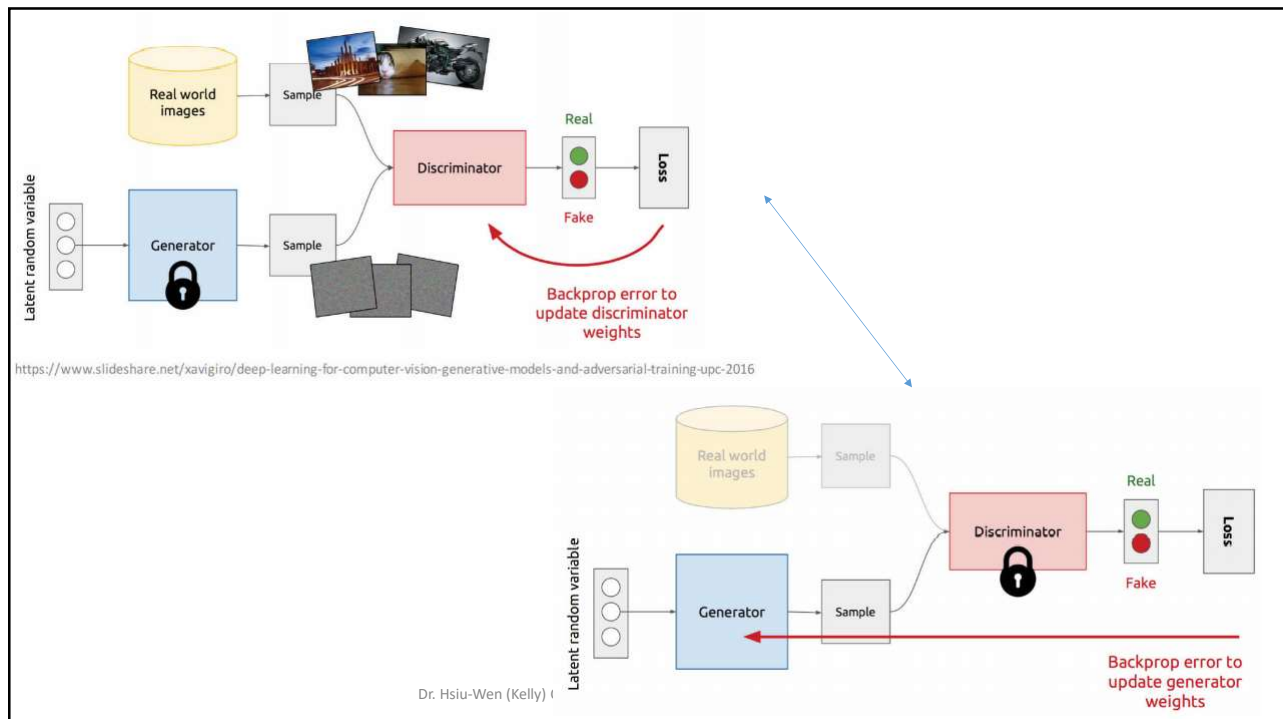
- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by ascending its stochastic gradient (improved objective):

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(z^{(i)})))$$


end for

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines ParisTech, PSL

33



- Deep Convolutional Generative Adversarial Networks



Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

35

Samples from the model look amazing!

Radford et al,
ICLR 2016




Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

36

Trajectory in latent space


- Walking in the latent space (z) can help us understand the landscape of it as well as to reason if the model has learned relevant and interesting representations




Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

37

Conditional GAN



- Training a conditional GAN to map edges→photo. The discriminator, D , learns to classify between fake (synthesized by the generator) and real {edge, photo} tuples. The generator, G , learns to fool the discriminator. Unlike an unconditional GAN, both the generator and discriminator observe the input edge map



Interactive demo: <https://affinelayer.com/pixsrv/>

Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." 2017.

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

38




"Video-to-Video Synthesis", NeurIPS'2018 [Nvidia+MIT]
Using Generative Adversarial Network (GAN)



Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

39

Source->Target domain transfer



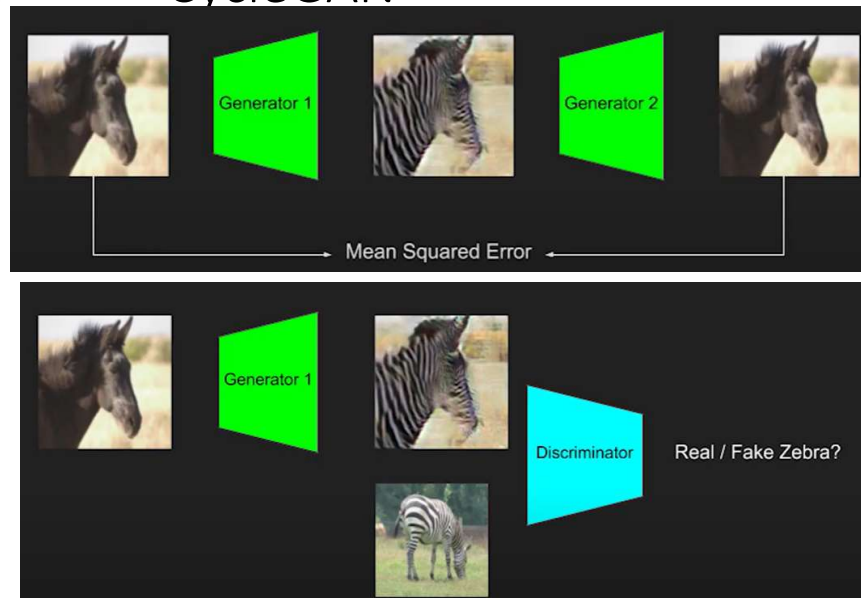
CycleGAN. Zhu et al. 2017.

- CycleGAN is proposed in 2017 (Jun-Yan Zhu, Taesung Park, et al.) to deal with task with Unpaired Image-to-Image Translation
- The architecture use two generators and on discriminator.
- The Objectives are
 - Ensure the translated image looks like zebra
 - This is trained using the GAN objective with the discriminator
 - Ensure the translated image still looks mostly like the original
 - This is trained using a reconstruction objective with the second generator
 - This is the novel cycle-consistency loss

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

40

PSL CycleGAN



41

"The GAN Zoo"

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AfiGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- iCGAN - Invertible Conditional GANs for image editing
- iD-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

<https://github.com/hindupuravinash/the-gan-zoo>

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines ParisTech, PSL

42



GANs short summary

- Don't work with an explicit density function
- Take game-theoretic approach: learn to generate from training distribution through 2-player game
- Pros:
 - Beautiful, state-of-the-art samples!
- Cons:
 - Trickier / more unstable to train
 - Can't solve inference queries such as $p(x)$, $p(z|x)$
- Active areas of research: -
 - Better loss functions, more stable training (Wasserstein GAN, LSGAN, many others)
 - Conditional GANs, GANs for all kinds of applications

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

43



Summary

- Intrinsically UNSUPERVISED
 - can be used on UNLABELLED DATA
- Impressive results in Image Retrieval
- DBN/DBM/VAE = Generative probabilistic models
- GAN = most promising generative model, with already many remarkable & exciting applications
- Strong potential for enhancement of datasets and for ultra-realistic synthetic data
- Interest for "creative" /artistic computing?

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

44



Any QUESTIONS ?

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

45