

Session 4

Recurrent Neural Network

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

1



PSL Acknowledgements

- The materials majorly derived from Prof. Fabien Moutarde. Some slides come from the online classes.
 - Fei-Fei Li + J.Johnson + S.Yeung: slides on "Recurrent Neural Networks" from the "Convolutional Neural Networks for Visual Recognition" course at Stanford

http://cs231n.stanford.edu/slides/2019/cs231n 2019 lecture10.pdf

• Yingyu Liang: slides on "Recurrent Neural Networks" from the "Deep Learning Basics" course at Princeton

https://www.cs.princeton.edu/courses/archive/spring16/cos495/slides/DL_lecture9_RNN.pdf

 Arun Mallya: slides "Introduction to RNNs" from the "Trends in Deep Learning and Recognition" course of Svetlana LAZEBNIK at University of Illinois at Urbana-Champaign

http://slazebni.cs.illinois.edu/spring17/lec02 rnn.pdf

 Tingwu Wang: slides on "Recurrent Neural Network" for a course at University of Toronto

https://www.cs.toronto.edu/%7Etingwuwang/rnn_tutorial.pdf

 Christopher Olah: online tutorial "Understanding LSTM Networks" https://colah.github.io/posts/2015-08-Understanding-LSTMs/

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

2



PSL* Introduction

- Recurrent neural networks have been an important focus of research and development during the 1990's. → It is much older than ConvNet!
- They are designed to learn sequential or time varying patterns.
- A recurrent net is a neural network with feedback (closed loop) connections [Fausett, 1994]. Examples include BAM, Hopfield, Boltzmann machine, and recurrent backpropagation nets [Hecht-Nielsen, 1990].
- A dynamic neural network can be defined as a neural networks that consists of interlayer feedback loops (i.e., from output layer to input layer) and intra-layer feedback loops (i.e., between different neurons within the same layer) or self-feedback loops.
- From the computational perspective, a dynamic neural network that contains the feedback loop that may provide more computational advantages than a static neural network, which contains only feed-forward architecture
- Applications: natural language processing (NLP), forecasting, signal processing and control require the treatment of dynamics associated with the unknown model.

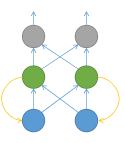
3

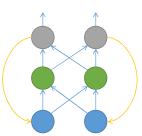
3



PSL™ Old style of RNN

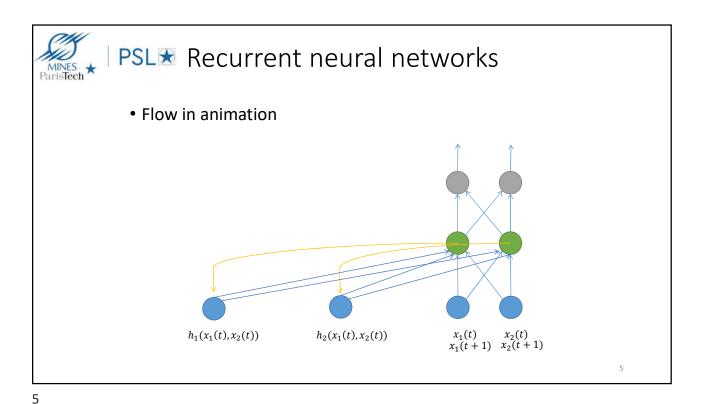
- Elman introduced feedback from the hidden layer to the context portion of the input layer.
 - This approach pays more attention to the sequence of input values.
- Jordan recurrent neural networks use feedback from the output layer to the context nodes of the input layer and give more emphasis to the sequence of output values.





However, these methods did not succeed in bigger data set due to the design of gradient flow

4



PSL

★ Flexibility • In some context of machine learning, we want to have flexibility of input and output one to one one to many many to one many to many many to many Image captioning Sentiment Video Language Image->seq of classification translation classification on words seq of words -> Seq of words -> frame level sentiment seq of words Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL



PSL Gradient flow

- Gradient flow is very important in network
- We already saw a lots in the last section
- Risk to have feed-back connection:
 - stability
 - Controllability
 - Observability

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

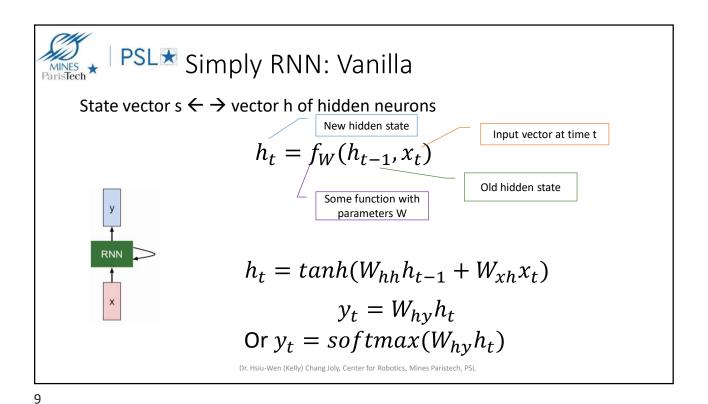
7

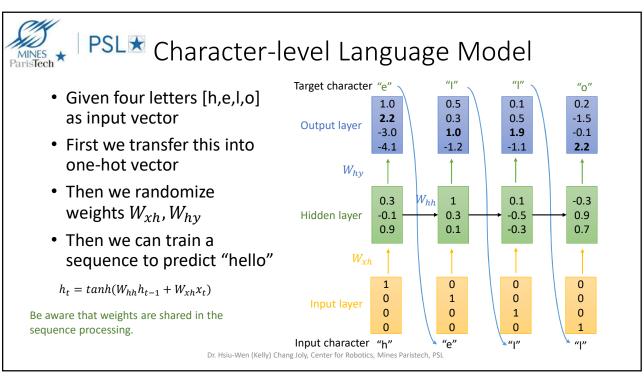


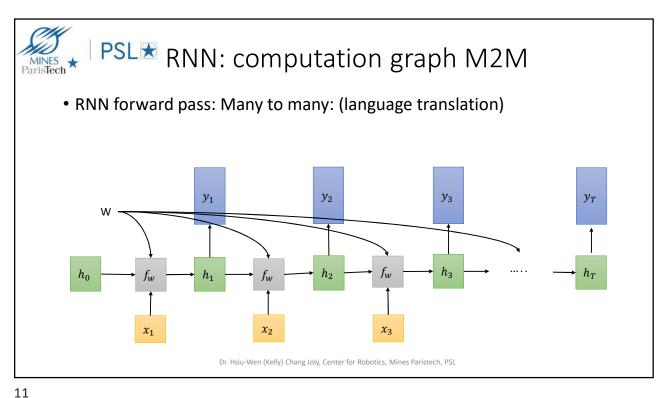
PSL★ Advantages of RNN

- The hidden state s of the RNN builds a kind of lossy summary of the past
- RNN totally adapted to processing SEQUENTIAL data (same computation formula applied at each time step, but modulated by the evolving "memory" contained in state s)
- Universality of RNNs: any function computable by a Turing Machine can be computed by a finite-size RNN (Siegelmann and Sontag, 1995)

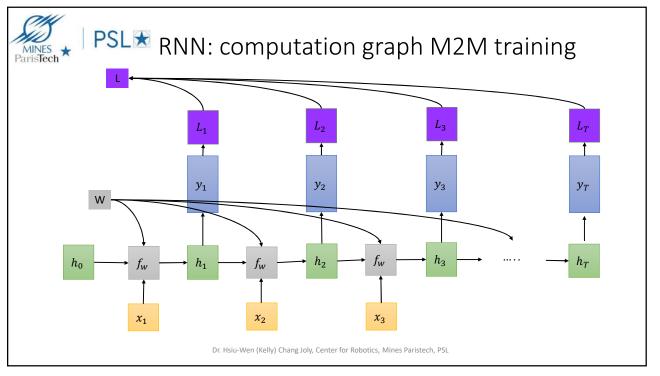
Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

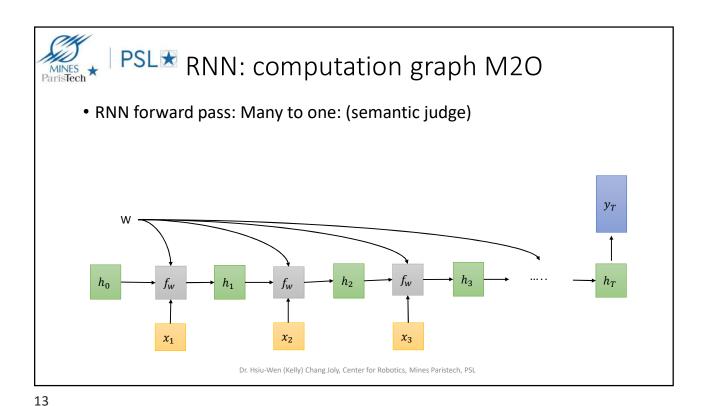






т 1

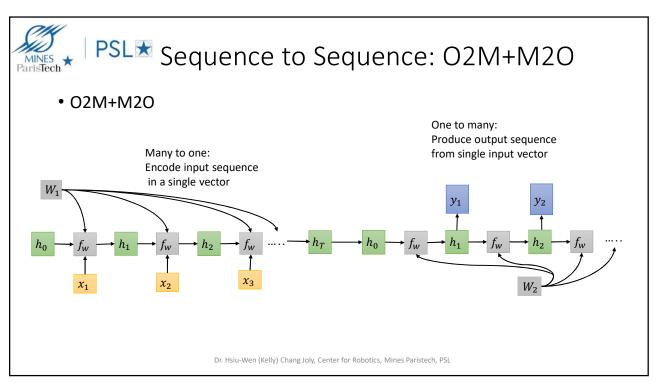




PSL \bigstar RNN: computation graph O2M

• RNN forward pass: One to many: (captioning) y_1 y_2 y_3 y_4 y_4 y_4 y_4 y_4 y_5 y_7 y_8 $y_$

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL





- Applying backpropagation in RNNs is called backpropagation through time [Werbos.1990].
- This procedure requires us to expand (or unroll) the computational graph of an RNN one time step at a time.
- The unrolled RNN is essentially a feedforward neural network with the special property that the same parameters are repeated throughout the unrolled network, appearing at each time step.
- Then, like in feedforward neural network, we apply the chain rule to backpropagate gradients through the unrolled net.
- The gradient with respect to each parameter must be summed across all places that the parameter occurs in the unrolle d net.
- Handling such weight tying should be familiar from our chapters on convolutional neural networks.

Werbos, P. J. (1990). Backpropagation through time: what it does and how to do it. *Proceedings of the IEEE*, 78(10), 1550–1560.



PSL★ Analysis of gradients

$$h_t = f(x_t, h_{t-1}, w_h)$$

o_t = $g(h_t, w_o)$

- where f and g are transformations of the hidden layer and the output layer, respectively.
- · Hence, we have a chain of values that depend on each other via recurrent computation

$$\{..., (x_{t-1}, h_{t-1}, o_{t-1}), (x_t, h_t, o_t), ...\}$$

The forward pass of this model is to loop through the (x_t, h_t, o_t) triples one time step at a time. The discrepancy between output o_t and the desired target y_t is then evaluated by an objective function across all the T time steps

$$L(x_1, ..., x_T, y_1, ..., y_T, w_h, w_o) = \frac{1}{T} \sum_{t=1}^{T} l(y_t, o_t)$$

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSI

17



PSL★ Analysis of gradients

 Here are the tricky way to derive the gradients regarding the parameters w_h of the objective function L.

$$\frac{\partial L}{\partial w_h} = \frac{1}{T} \sum_{t=1_T}^T \frac{\partial l(y_t, o_t)}{\partial w_h}$$

$$= \frac{1}{T} \sum_{t=1}^T \frac{\partial l(y_t, o_t)}{\partial o_t} \frac{\partial (o_t)}{\partial w_h} = \frac{1}{T} \sum_{t=1}^T \frac{\partial l(y_t, o_t)}{\partial o_t} \frac{\partial g(h_t, w_o)}{\partial h_t} \frac{\partial h_t}{\partial w_h}$$

• The third term in this equation is tricky because the computation of h_t depends on both h_{t-1} , w_h where computation of h_{t-1} also depends on w_h . Applying total differential of $df(x,y) = f_x dx + f_y dy$, we have: $\frac{\partial h_t}{\partial w_h} = \frac{\partial f(x_t,h_{t-1},w_h)}{\partial w_h} + \frac{\partial f(x_t,h_{t-1},w_h)}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial w_h}$

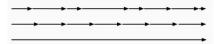
$$\frac{\partial h_t}{\partial w_h} = \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial w_h} + \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial w_h}$$



PSL Analysis of gradients

$$\frac{\partial h_t}{\partial w_h} = \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial w_h} + \sum_{i=1}^{t-1} \left(\prod_{j=i+1}^{t} \frac{\partial f(x_j, h_{j-1}, w_h)}{\partial h_{j-1}} \right) \frac{\partial f(x_i, h_{i-1}, w_h)}{\partial w_h}$$

- While we can use the chain rule to compute $\frac{\partial h_t}{\partial w_h}$ recursively, this chain can get very long whenever t is large.
 - Full computation: very slow and gradients can blow up, since subtle changes in the initial conditions can potentially affect the outcome a lot
 - Truncating time steps [Jaeger, 2002]
 - Randomized Truncation [Tallec and Ollivier, 2017]



Jaeger, H. (2002). Tutorial on training recurrent neural networks, covering BPPT, RTRL, EKF and the echo state network approach. Vol. 5. GMD-Forschungszentrum Informationstechnik Bonn

Tallec, C., & Ollivier, Y. (2017). Unbiasing truncated backpropagation through time. arXiv preprint arXiv:1705.08209.

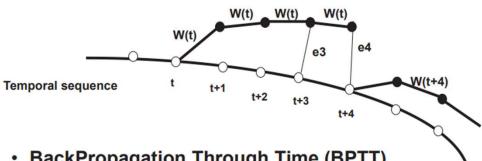
Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

19

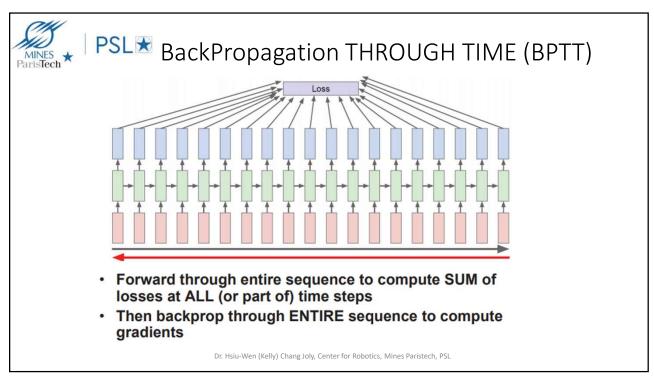


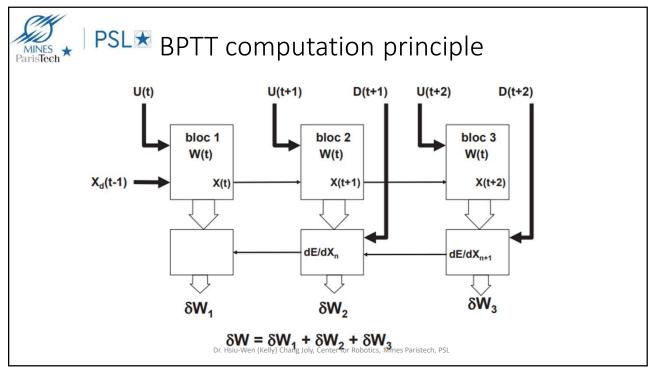
PSL★ RNN training

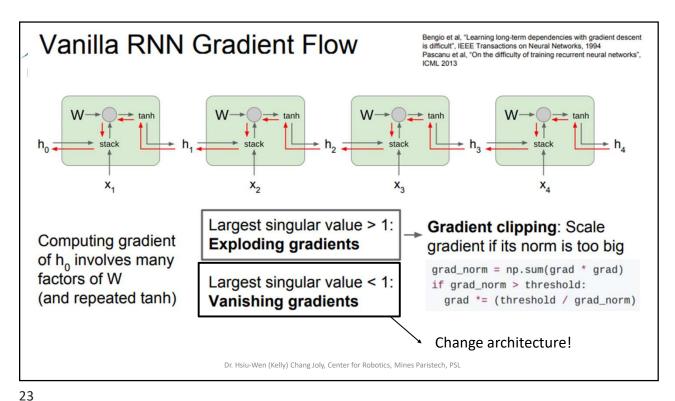
Horizon $N_t = 4$



- BackPropagation Through Time (BPTT) gradients update for a whole sequence
- or Real Time Recurrent Learning (RTRL) gradients update for each frame in a sequence





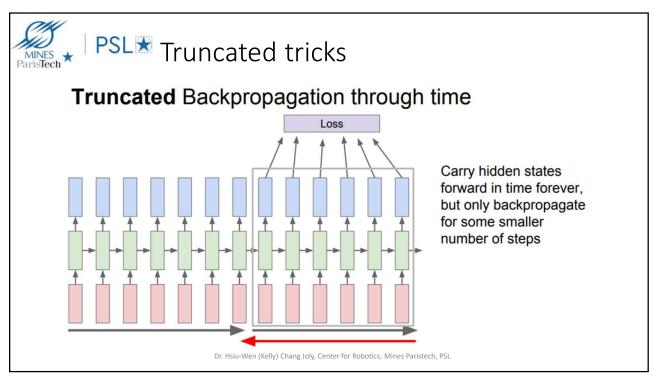


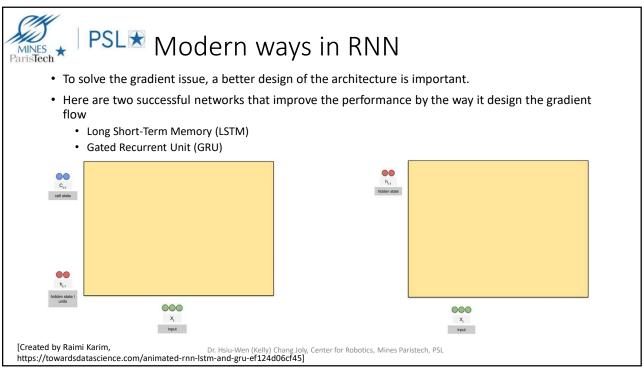
__



- If eigenvalues of Jacobian matrix >1, then gradients tend to EXPLODE
 → Learning will never converge.
- Conversely, if eigenvalues of Jacobian matrix < 1, then gradients tend to VANISH
 - → Error signals can only affect small time lags
 - →short-term memory.
- Possible solutions for exploding gradient: CLIPPING trick (limited values in an array, see numpy.clip), truncated.
- Possible solutions for vanishing gradient:
 - · use ReLU instead of tanh
 - change what is inside the RNN!

Recommended code to read for better understand this slide: https://gist.github.com/karpathy/d4dee566867f8291f086







PSL™ Long short-term memory

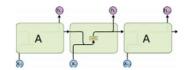
- The term "long short-term memory" comes from the following intuition. Simple recurrent neural networks have long-term memory in the form of weights. The weights change slowly during training, encoding general knowledge about the data. They also have short-term memory in the form of ephemeral activations, which pass from each node to successive nodes. The LSTM model introduces an intermediate type of storage via the memory cell. A memory cell is a composite unit, built from simpler nodes in a specific connectivity pattern, with the novel inclusion of multiplicative nodes
- Gated memory cell is equipped with an internal state and a number of multiplicative gates that determine
 - a given input should impact the internal state (the *input gate*): $i \in [0,1]$
 - Weather the internal state should be flushed to 0 (the forget gate): $f \in [0,1]$
 - Weather the internal state of a given neuron should be allowed to impact the cell's output (the *output* gate): $o \in [0,1]$

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

27



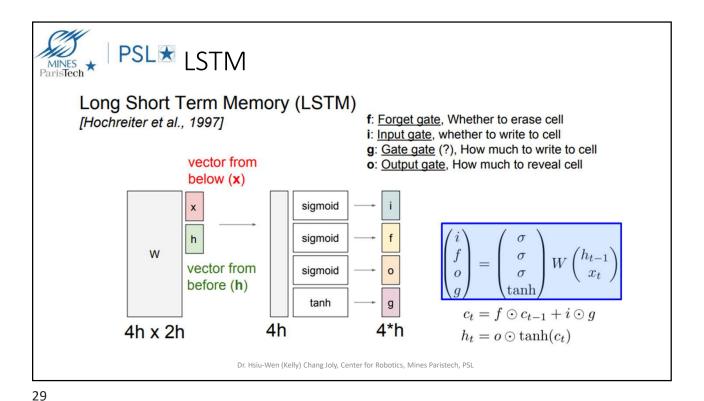
Problem of standard RNNs = no actual LONG-TERM memory



LSTM = RNN variant for solving this issue (proposed by Hochreiter & Schmidhuber in 1997)

[Figures from https://colah.github.io/posts/2015-08-Understanding-LSTMs/]

 Key idea = use "gates" that modulate respective influences of input and memory



Gate = $\underbrace{pointwise}_{\text{ParisTech}}$ multiplication by σ in]0;1[\Rightarrow modulate between "let nothing through"
 and "let everything through"

• FORGET gate $f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$ • INPUT gate $i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$ $\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$ \Rightarrow next state = mix between

pure memory or pure new $C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$ [Figures from https://colah.github.io/posts/2015-08-Understanding-LSTMs/]

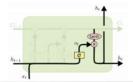


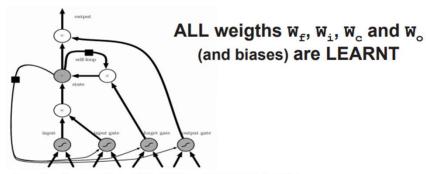
PSL★ LSTM summary

OUTPUT gate

$$o_t = \sigma \left(W_o \left[h_{t-1}, x_t \right] + b_o \right)$$

$$h_t = o_t * \tanh \left(C_t \right)$$





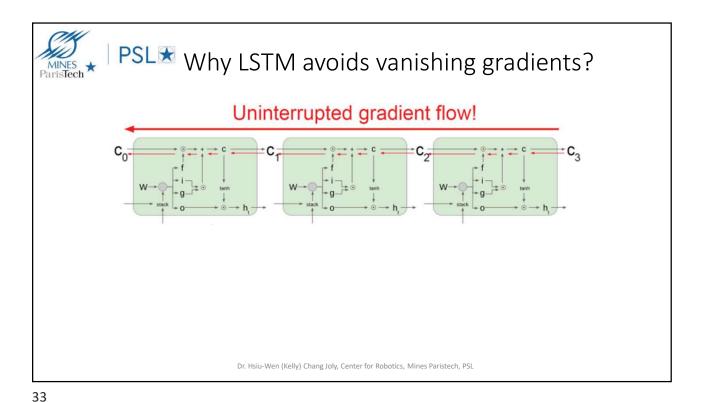
[Figure from Deep Learning book by I. Goodfellow, Y. Bengio & A. Courville]

31



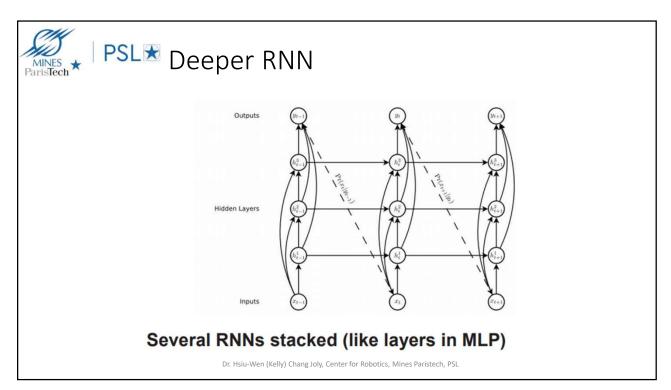
PSL Practice

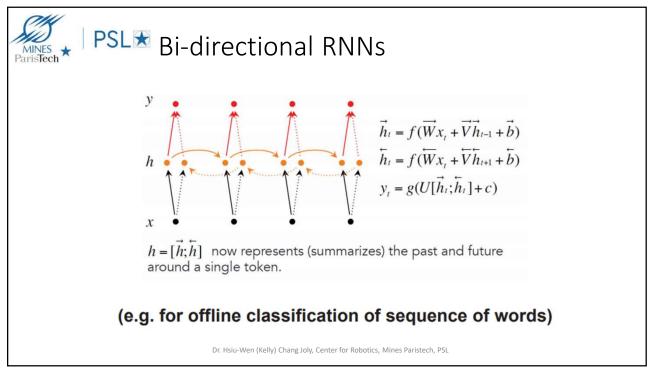
- Given input $x_t \in R^4$ and the target y_t at time step t. The hidden state $h_t \in R^2$ and the output $o_t \in R^1$.
 - What is the size of the weight when vanilla RNN is applied?
 - What is the size of the weights when LSTM is applied?

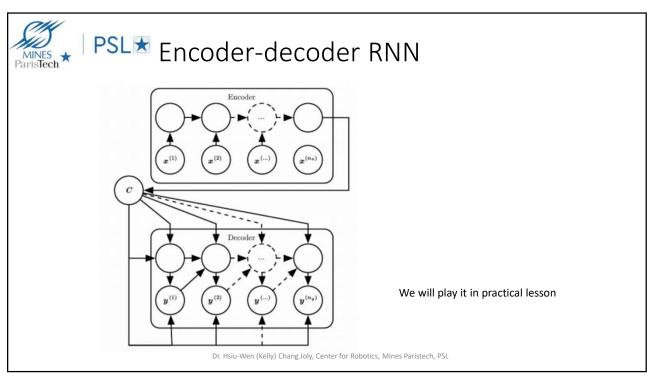


Simplified variant of LSTM, with only 2 gates:
 a RESET gate & an UPDATE gate
 (proposed by Cho, et al. in 2014)

GRU [Learning phrase representations using rnn encoder-decoder for statistical machine translation, Cho et al. 2014] $t_t = \sigma(W_{xx}t_t + W_{hx}h_{t-1} + b_r)$ $t_t = \sigma(W_{xx}t_t + W_{hx}h_{t-1} + b_z)$ $t_t = \tau(W_{xx}t_t + W_{hx}h_{t-1} + b_z)$

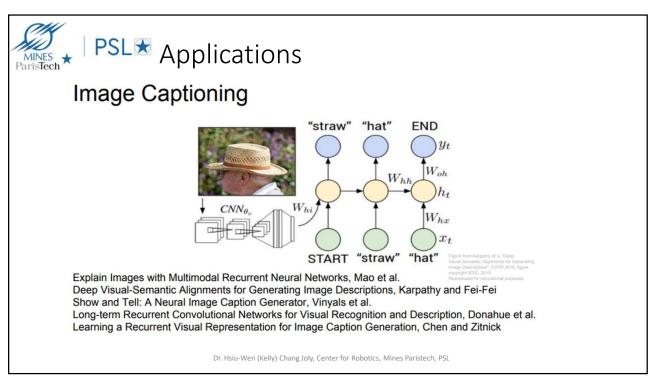


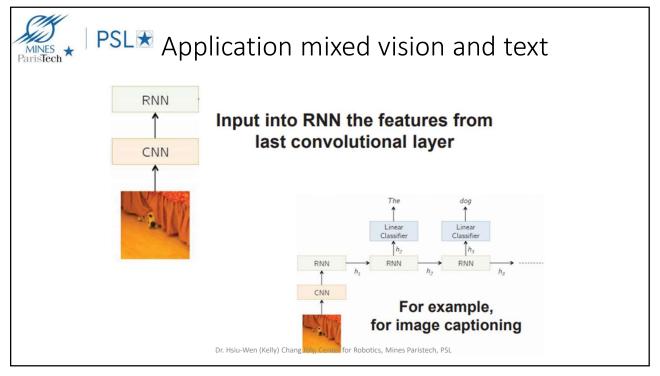


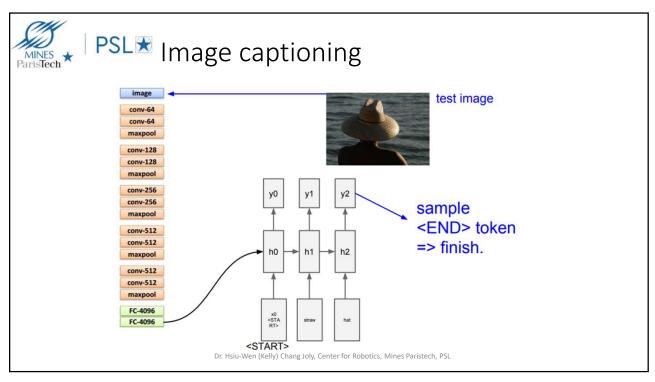




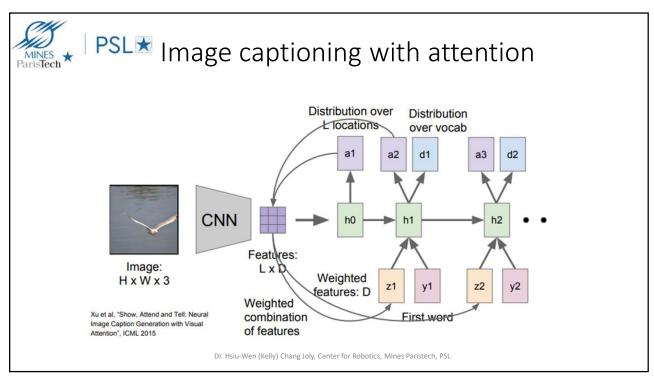
- [LSTM: A Search Space Odyssey, Greff et al., 2015] :
 - they play around the LSTM equations, swap out the non linearities at one point, do we need tanh, this paper made a lot of experiences in playing around different design
 - Conclusion is there is no significant difference.
- [An Empirical Exploration of Recurrent Network Architectures, Jozefowicz et al., 2015]:
 - Search over very large number of random RNN architecture, randomly permute these equations to see if there is a better one
 - Conclusion: No significant improvement with one specific version















PSL★ Visual Question Answering



- What endangered animal is featured on the truck?
- A bald eagle. A: A sparrow.
- A: A humming bird.
- A: A raven.



- Q: Where will the driver go if turning right?
- A: Onto 24 3/4 Rd.
- A: Onto 25 3/4 Rd. A: Onto 23 3/4 Rd.
- A: Onto Main Street.



- Q: When was the picture taken?
- A: During a wedding. A: During a bar mitzvah.
- A: During a funeral.
- A: During a Sunday church



- Q: Who is under the umbrella?
- A: Two women.
- A: An old man. A: A husband and a wife.

Agrawal et al, "VQA: Visual Question Answering", ICCV 2015
Zhu et al, "Visual 7W: Grounded Question Answering in Images", CVPR 2016
Figure from Zhu et al, copyright IEEE 2016. Reproduced for educational purposes.

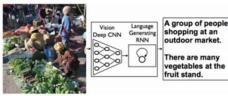
Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

45



Wherever data is intrinsicly SEQUENTIAL

- · Speech recognition
- Natural Language Processing (NLP)
 - Machine-Translation
 - Image caption generator



- Gesture recognition
- Music generation
- Potentially any kind of time-series!!

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL

Check GPT-3 that can be used in a large of different applications



- RNNs allow a lot of flexibility in architecture design
- Vanilla RNNs are simple but don't work very well
- Common to use LSTM or GRU: their additive interactions improve gradient flow
- Backward flow of gradients in RNN can explode or vanish
- Exploding is controlled with gradient clipping
- Vanishing is controlled with additive interactions (LSTM)
- Better/simpler architectures are a hot topic of current research
- Better understanding (both theoretical and empirical) is needed.

Dr. Hsiu-Wen (Kelly) Chang Joly, Center for Robotics, Mines Paristech, PSL