

# Session 6

Deep Generative model

Mini-project topics

# Acknowledgements

- The materials majorly derived

[http://cs231n.stanford.edu/slides/2017/cs231n\\_2017\\_lecture14.pdf](http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture14.pdf)

- OpenAI spinningup

[https://spinningup.openai.com/en/latest/spinningup/rl\\_intro.html](https://spinningup.openai.com/en/latest/spinningup/rl_intro.html)

- David Silver

[https://www.davidsilver.uk/wp-content/uploads/2020/03/intro\\_RL.pdf](https://www.davidsilver.uk/wp-content/uploads/2020/03/intro_RL.pdf)

- Introduction
- [Fully visible belief network](#)
- [Boltzmann machine/RBM/DBM](#)
- [Autoencoder](#)
- [GAN](#)

# Supervised vs Unsupervised

## Supervised Learning

**Data:**  $(x, y)$

$x$  is data,  $y$  is label

**Goal:** Learn a *function* to map  $x \rightarrow y$

**Examples:** Classification, regression, object detection, semantic segmentation, image captioning, etc.

## Unsupervised Learning

Training data is cheap

**Data:**  $x$

Just data, no labels!

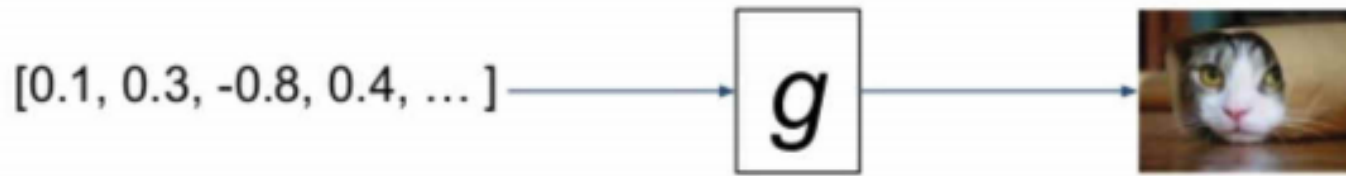
Holy grail: Solve  
unsupervised learning  
 $\Rightarrow$  understand structure  
of visual world

**Goal:** Learn some underlying  
hidden *structure* of the data

**Examples:** Clustering, dimensionality reduction, feature learning, density estimation, etc.

# Unsupervised Learning

- Examples:
  - Dimension reduction : PCA
  - Clustering: k-means
  - Density estimation
  - Feature learning
- General framework:
  - Find deterministic function  $f$ :  $z=f(x)$ ,  $x$  is data,  $z$  is the latent



## Unsupervised learning vs. Generative model

- $z = f(x)$  vs.  $x = g(z)$
- $P(z|x)$  vs.  $P(x|z)$
- **Encoder** vs. **Decoder ( Generator )**
  - $P(x, z)$  needed. ( cf :  $P(y|x)$  in supervised learning )
    - $P(z|x) = P(x, z) / P(x)$
    - $P(x|z) = P(x, z) / P(z) \rightarrow P(z)$  is given. ( prior )



# Why generative models?

- Realistic samples for artwork, super-resolution, colorization, etc.



- Generative models of time-series data can be used for simulation and planning (reinforcement learning applications!)
- Training generative models can also enable inference of latent representations that can be useful as general features

# Taxonomy of Generative Models

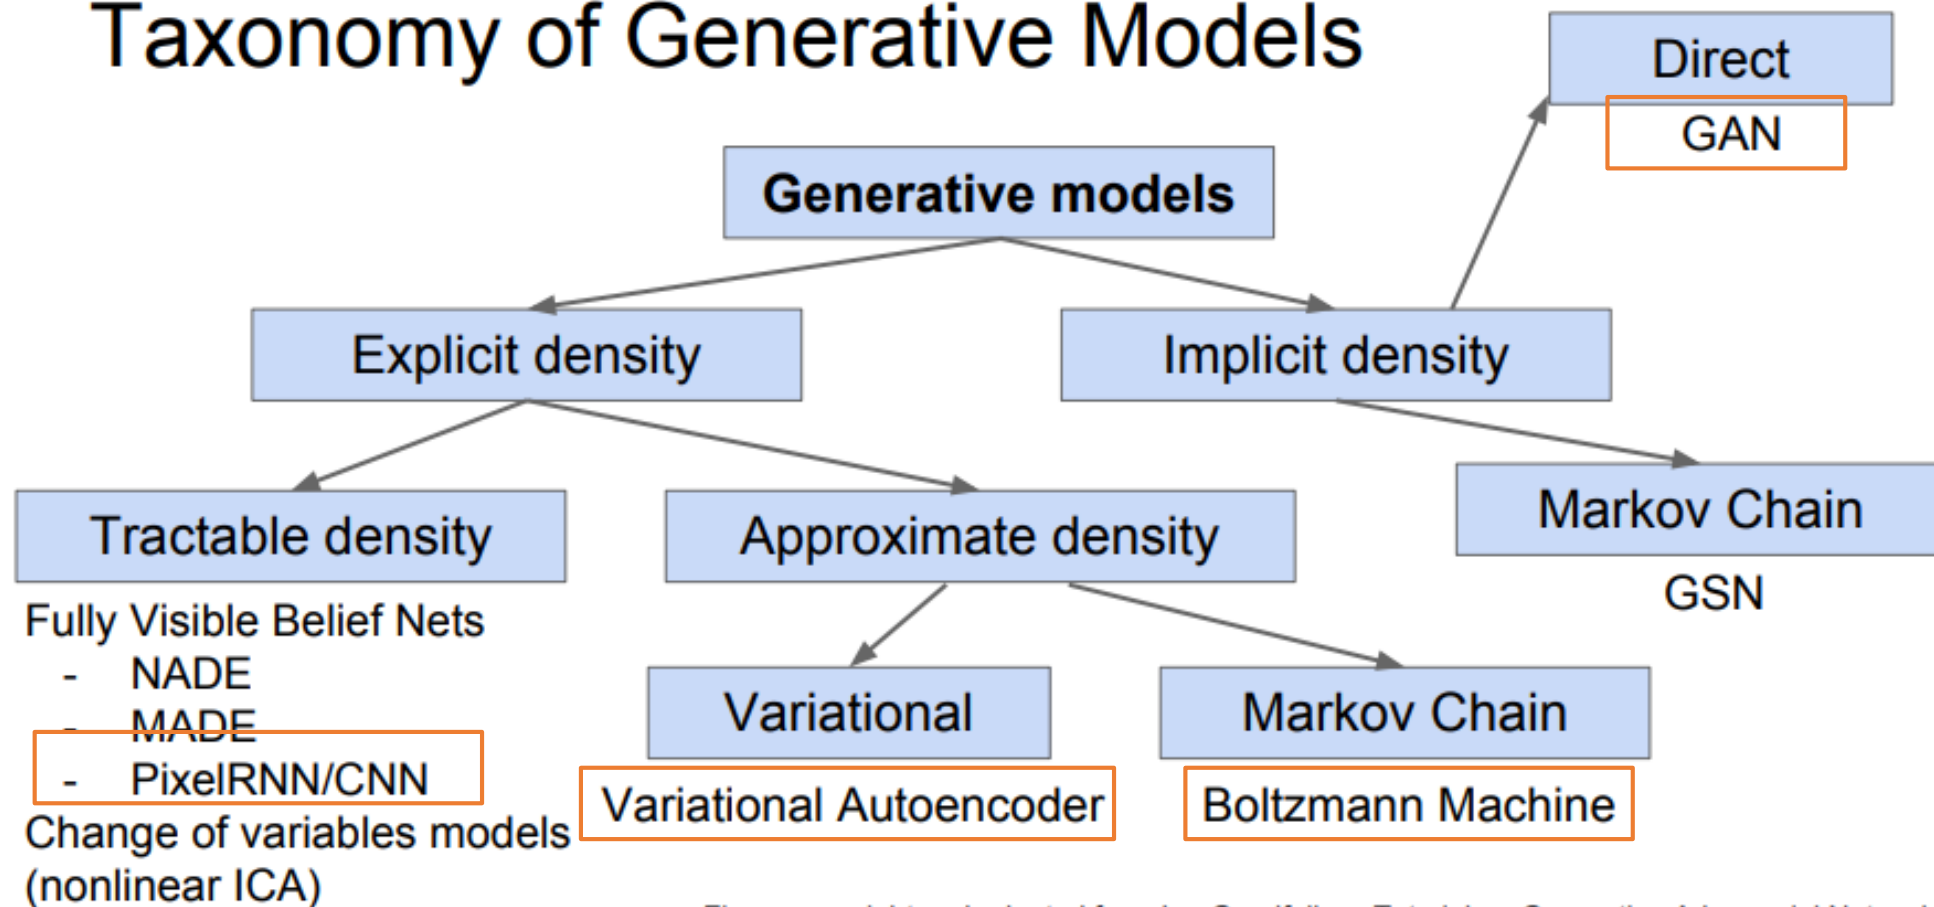


Figure copyright and adapted from Ian Goodfellow, Tutorial on Generative Adversarial Networks, 2017.



# Fully visible belief network

## Explicit density model

Use chain rule to decompose likelihood of an image  $x$  into product of 1-d distributions:

$$p(x) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

Likelihood of  
image x

Probability of i'th pixel value  
given all previous pixels

Will need to define ordering of “previous pixels”

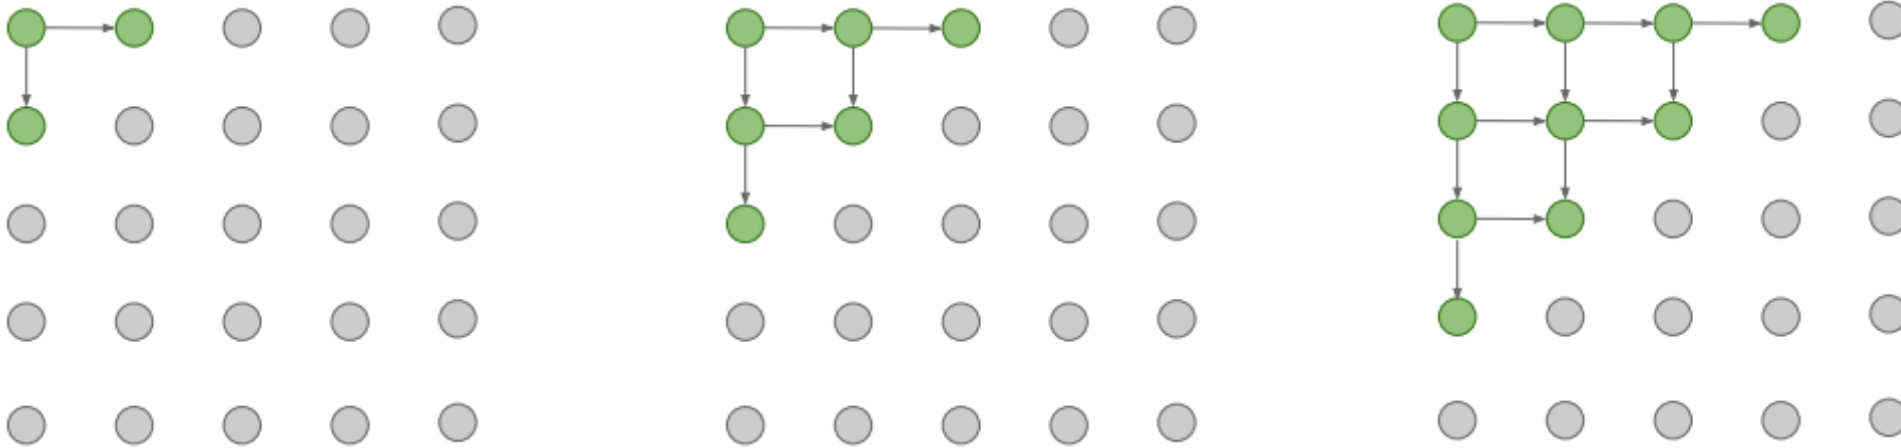
Complex distribution over pixel values => Express using a neural network!

Then maximize likelihood of training data

## Drawback: sequential generation is slow!

# PixelRNN [van der Oord et al. 2016]

- Generate image pixels starting from corner
- Dependency on previous pixels modeled using an RNN (LSTM)



# PixelCNN [van der Oord et al. 2016]

Still generate image pixels starting from corner

Dependency on previous pixels now modeled using a CNN over context region

Training is faster than PixelRNN  
(can parallelize convolutions since context region values known from training images)

Generation must still proceed sequentially  
=> still slow

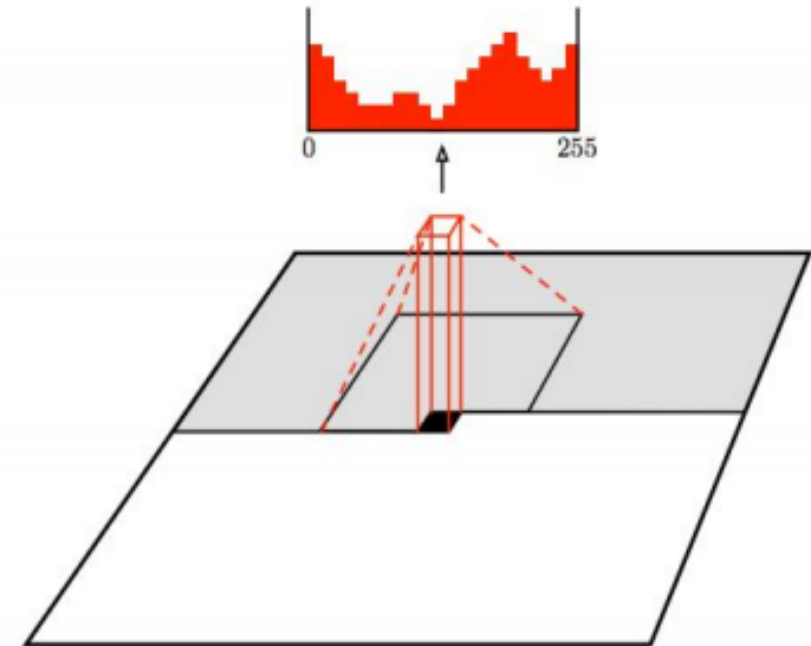


Figure copyright van der Oord et al., 2016. Reproduced with permission.

# PixelRNN and PixelCNN

## Pros:

- Can explicitly compute likelihood  $p(x)$
- Explicit likelihood of training data gives good evaluation metric
- Good samples

## Con:

- Sequential generation => slow

## Improving PixelCNN performance

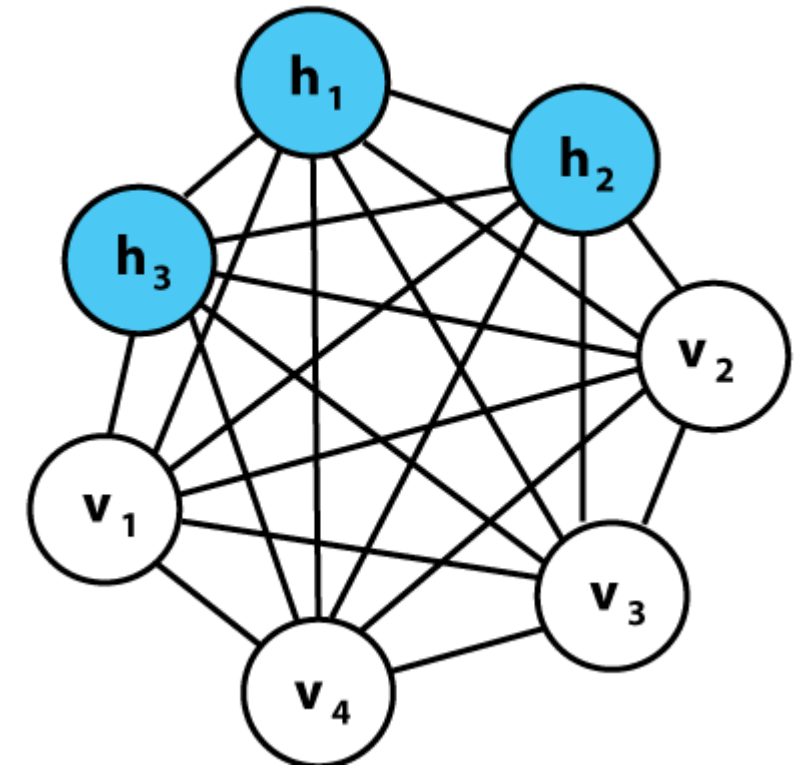
- Gated convolutional layers
- Short-cut connections
- Discretized logistic loss
- Multi-scale
- Training tricks
- Etc...

## See

- Van der Oord et al. NIPS 2016
- Salimans et al. 2017 (PixelCNN++)

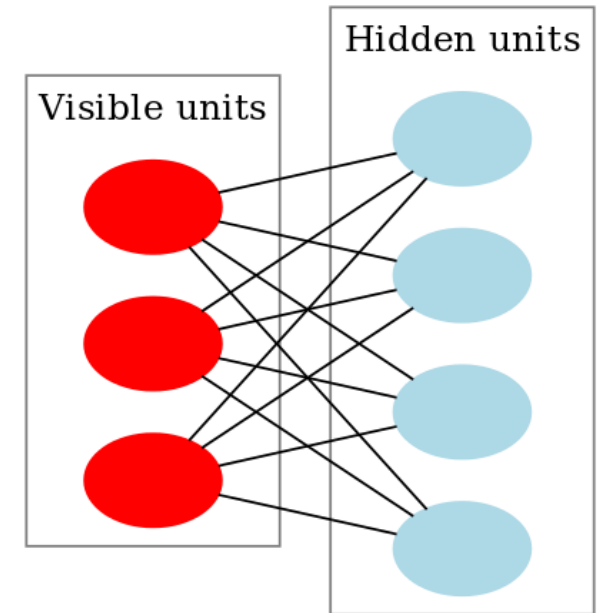
# What is Boltzmann machine?

- Boltzmann machines are stochastic and generative neural networks capable of learning internal representations
- They are able to represent (given sufficient time) and solve difficult combinatorial problems.
- They were invented in 1985 by Geoffrey Hinton
- non-deterministic (or stochastic) generative Deep Learning models with only two types of nodes: visible ( $v$ ) and hidden ( $h$ )
- Unlike classical neural networks
  - No output nodes
  - Connection between input nodes ( $v$ )
- This allows them to share information among themselves and self-generate subsequent data



# Restricted Boltzmann machine (RBM)

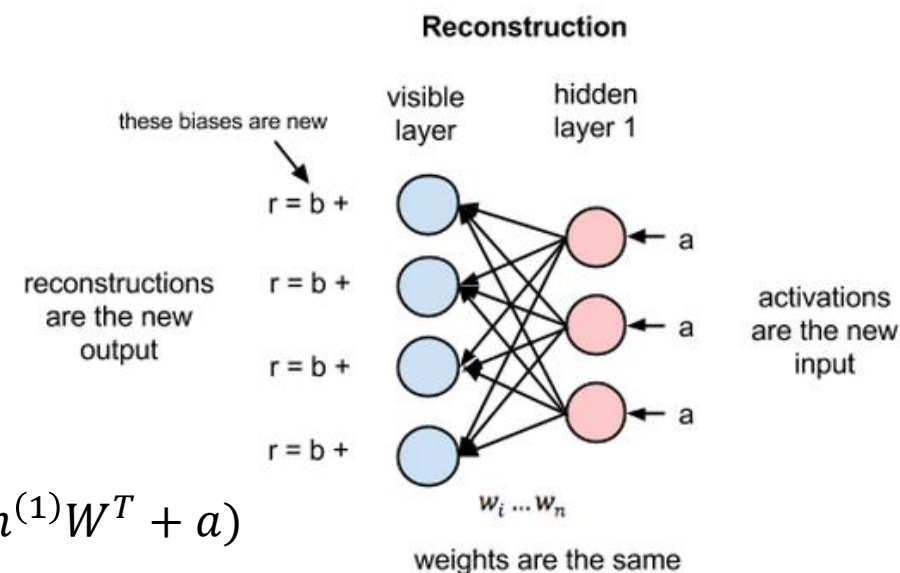
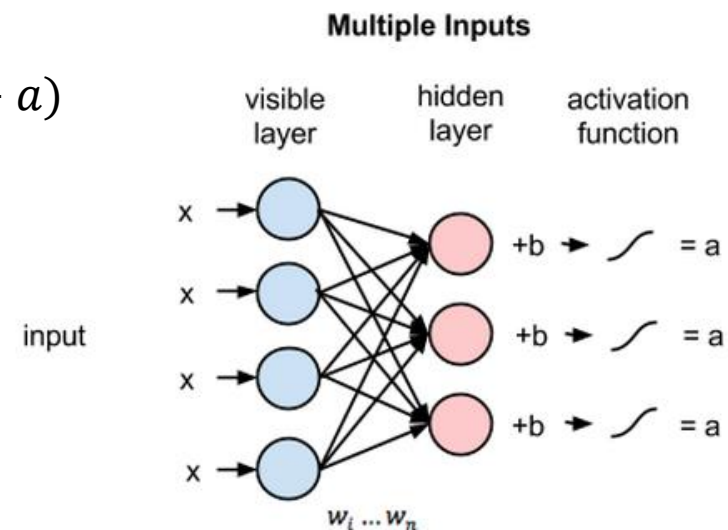
- RBMs are a two-layered artificial neural network with generative capabilities.
- They have the ability to learn a probability distribution over its set of input.
- It can be used for dimensionality reduction, classification, regression, collaborative filtering, feature learning, and topic modeling.
- RBMs are a special class of Boltzmann Machines and they are restricted in terms of the connections between the visible and the hidden units.
- Every node in the visible layer is connected to every node in the hidden layer but no two nodes in the same group are connected to each other.
- This restriction allows for more efficient training algorithms than what is available for the general class of Boltzmann machines, in particular, the gradient-based contrastive divergence algorithm.





- RBM is a Stochastic Neural Network which means that each neuron will have some random behavior when activated.
- There are two other layers of bias units (hidden bias and visible bias) in an RBM.
- The hidden bias RBM produce the activation on the forward pass and the visible bias helps RBM to reconstruct the input during a backward pass.
- The reconstructed input is always different from the actual input as there are no connections among the visible units and therefore, no way of transferring information among themselves.

$$h^1 = S(v^{(0)T}W + a)$$

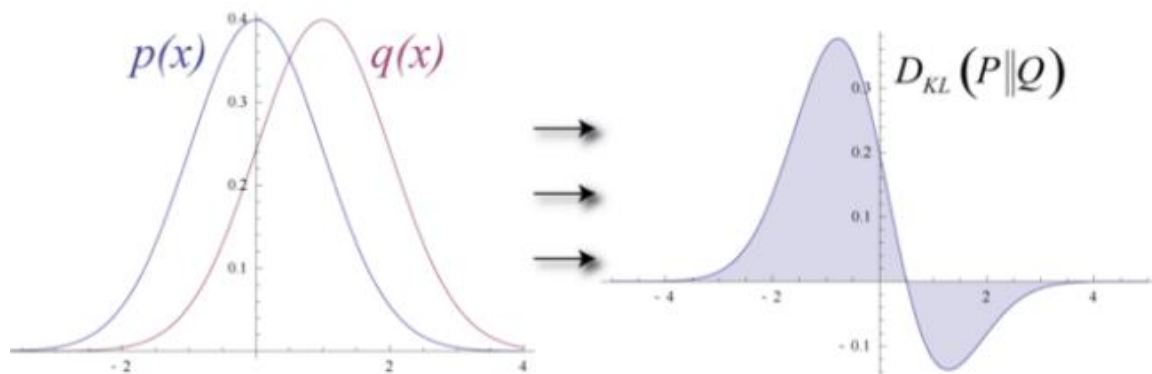


$$v^1 = S(h^{(1)}W^T + a)$$

# Learning process of RBM

- Consider the difference  $v(0)-v(10)$  as reconstruction error. The weights are adjusted in each iteration so as to minimize this error
- In forward pass:  $p(h^{(1)}|v^{(0)}; W)$
- In backward pass:  $p(v^{(1)}|h^{(1)}; W)$
- It is a joint distribution:  $p(v, h)$
- Assume that we have two normal distributions, one from the input data (denoted by  $p(x)$ ) and one from the reconstructed input approximation (denoted by  $q(x)$ ). The difference between these two distributions is our error in the graphical sense and our goal is to minimize it → Kullback-Leibler divergence (KL-divergence)
- KL-divergence measures the non-overlapping areas under the two graphs and the RBM's optimization algorithm tries to minimize this difference by changing the weights so that the reconstruction closely resembles the input.

# Contrastive divergence



Original Gaussian PDF's

KL Area to be Integrated

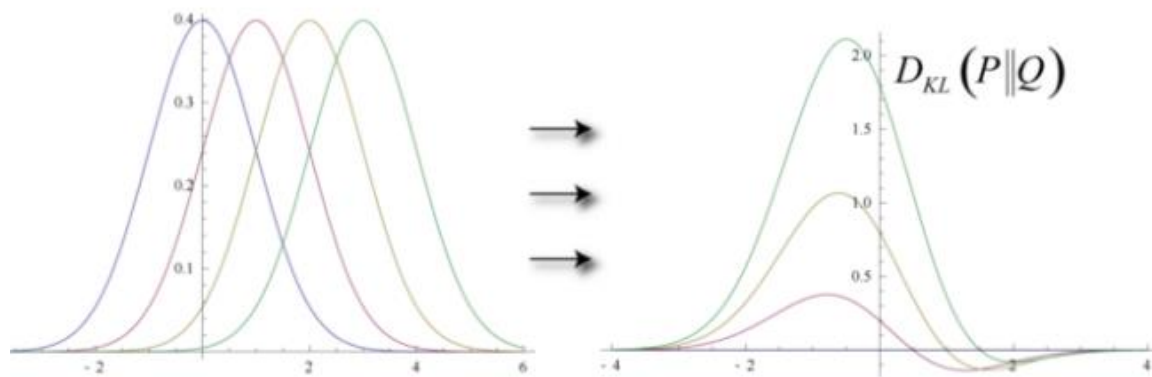


Image by [Mundhenk](#) on [Wikimedia](#)

Boltzmann Machines (and RBMs) are Energy-based models and a joint configuration,  $(\mathbf{v}, \mathbf{h})$  of the visible and hidden units has an energy given by:

$$E(\mathbf{v}, \mathbf{h}) = - \sum_{i \in \text{visible}} a_i v_i - \sum_{j \in \text{hidden}} b_j h_j - \sum_{i,j} v_i h_j w_{ij}$$

The probability that the network assigns to a visible vector,  $\mathbf{v}$ , is given by summing over all possible hidden vectors

$$p(\mathbf{v}) = \frac{\sum_{\mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{v}, \mathbf{h}} e^{-E(\mathbf{v}, \mathbf{h})}} \quad \frac{\partial \log p(\mathbf{v})}{\partial w_{ij}} = \langle v_i h_j \rangle_{\text{data}} - \langle v_i h_j \rangle_{\text{model}}$$

$$\Delta w_{ij} = \alpha (\underbrace{\langle v_i h_j \rangle_{\text{data}}}_{\text{Expectation}} - \underbrace{\langle v_i h_j \rangle_{\text{model}}}_{\text{reconstruct}})$$

Learning rate

# Contrastive divergence

- The learning rule is much more closely approximating the gradient of another objective function called the **Contrastive Divergence** which is the difference between two Kullback-Liebler divergences

---

## Algorithm 1. $k$ -step contrastive divergence

---

**Input:** RBM  $(V_1, \dots, V_m, H_1, \dots, H_n)$ , training batch  $S$

**Output:** gradient approximation  $\Delta w_{ij}$ ,  $\Delta b_j$  and  $\Delta c_i$  for  $i = 1, \dots, n$ ,  
 $j = 1, \dots, m$

```

1 init  $\Delta w_{ij} = \Delta b_j = \Delta c_i = 0$  for  $i = 1, \dots, n, j = 1, \dots, m$ 
2 forall the  $v \in S$  do
3    $v^{(0)} \leftarrow v$ 
4   for  $t = 0, \dots, k - 1$  do
5     for  $i = 1, \dots, n$  do sample  $h_i^{(t)} \sim p(h_i | v^{(t)})$ 
6     for  $j = 1, \dots, m$  do sample  $v_j^{(t+1)} \sim p(v_j | h^{(t)})$ 
7   for  $i = 1, \dots, n, j = 1, \dots, m$  do
8      $\Delta w_{ij} \leftarrow \Delta w_{ij} + p(H_i = 1 | v^{(0)}) \cdot v_j^{(0)} - p(H_i = 1 | v^{(k)}) \cdot v_j^{(k)}$ 
9      $\Delta b_j \leftarrow \Delta b_j + v_j^{(0)} - v_j^{(k)}$ 
10     $\Delta c_i \leftarrow \Delta c_i + p(H_i = 1 | v^{(0)}) - p(H_i = 1 | v^{(k)})$ 

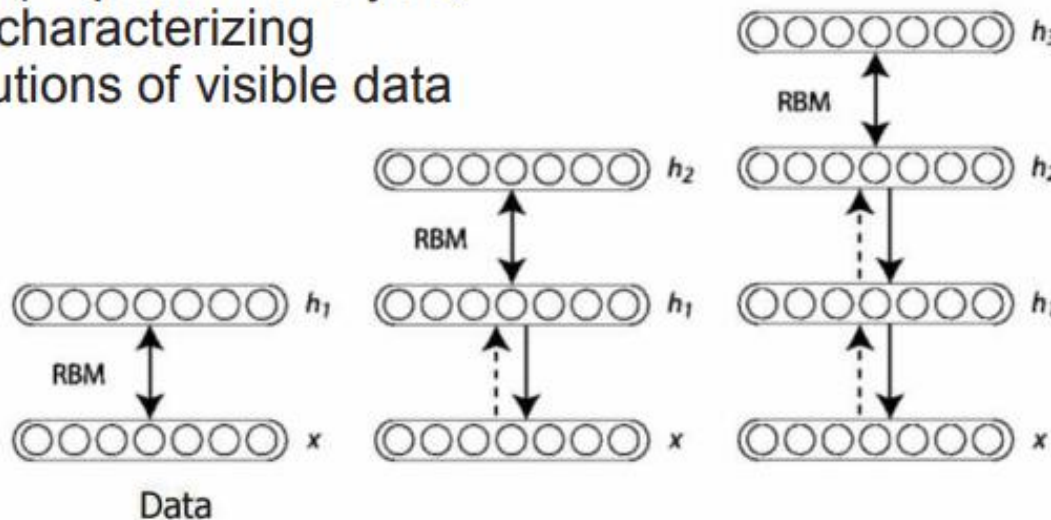
```

---

# Deep belief network (DBM)

- Multiple RBMs can be stacked and can be fine-tuned through the process of gradient descent and back-propagation.
- One of first Deep-Learning models
- Proposed by G. Hinton in 2006
- Generative probabilistic model (mostly UNSUPERVISED)

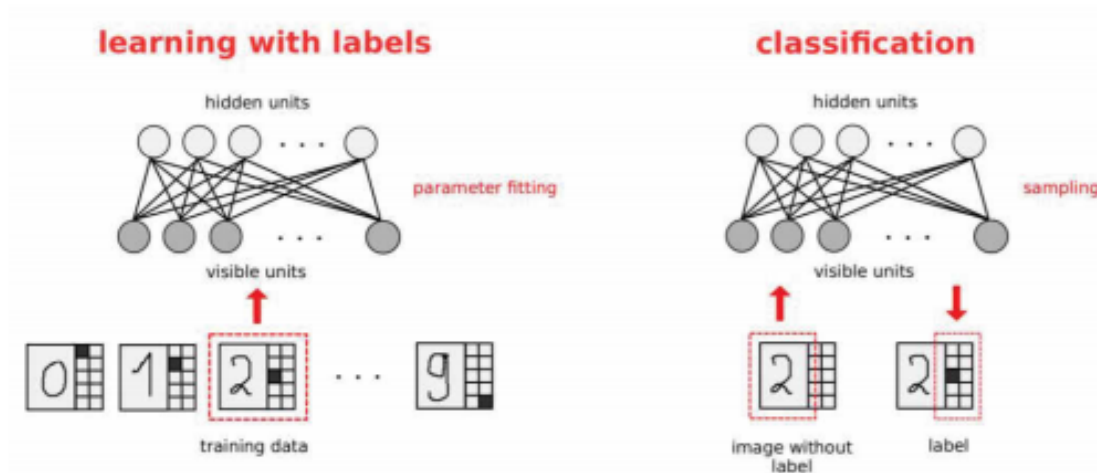
For capturing high-order *correlations* of observed/visible data ( $\rightarrow$  pattern analysis, or synthesis); and/or characterizing *joint* statistical distributions of visible data





# Use of trained RBM

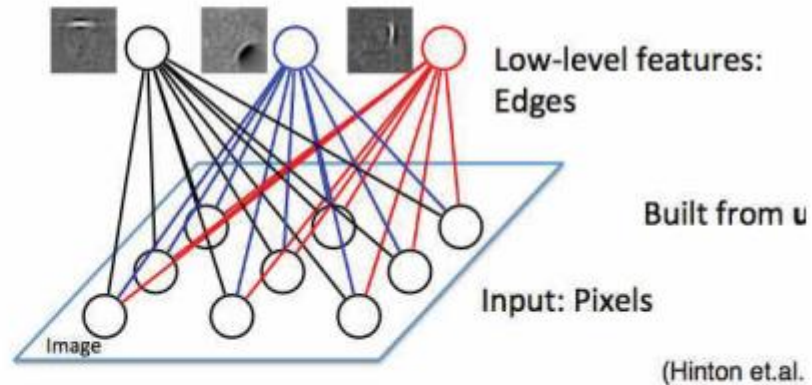
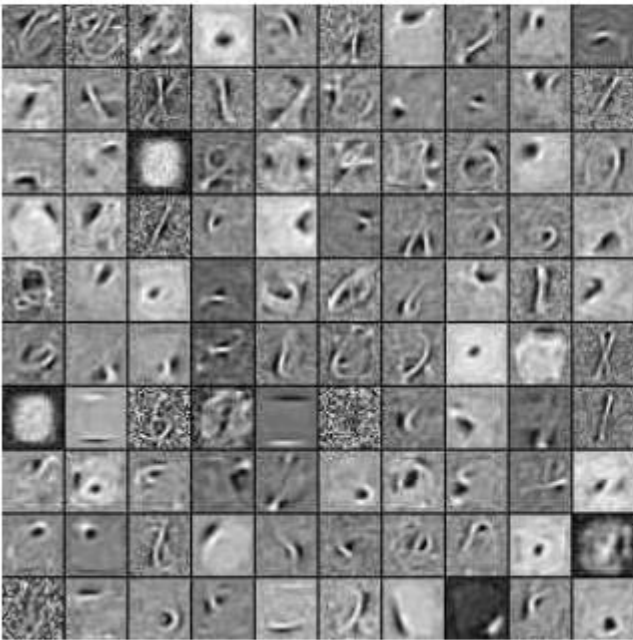
- Input data "completion" : set some  $v_i$  then compute  $h$ , and generate compatible full samples
- Generating representative samples
- Classification if trained with inputs=data+label





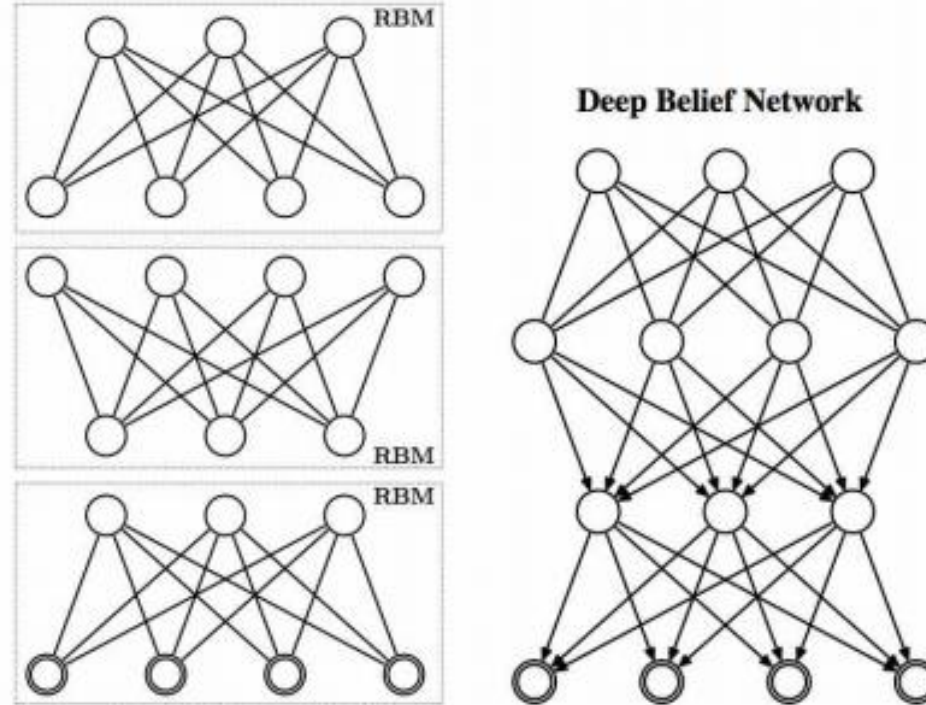
# Interpretation of trained RBM hidden layer

- Look at weights of hidden nodes → low-level features



This is why people are inspired to stack the RBMs to get more “abstract” features.

## Greedy learning of successive layers




---

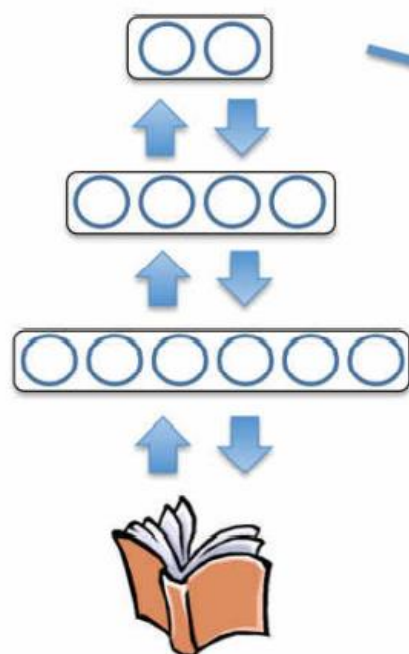
**Algorithm 1** Recursive Greedy Learning Procedure for the DBN.

---

- 1: Fit parameters  $W^1$  of the 1<sup>st</sup> layer RBM to data.
  - 2: Freeze the parameter vector  $W^1$  and use samples  $h^1$  from  $Q(h^1|v) = P(h^1|v, W^1)$  as the data for training the next layer of binary features with an RBM.
  - 3: Freeze the parameters  $W^2$  that define the 2<sup>nd</sup> layer of features and use the samples  $h^2$  from  $Q(h^2|h^1) = P(h^2|h^1, W^2)$  as the data for training the 3<sup>rd</sup> layer of binary features.
  - 4: Proceed recursively for the next layers.
-

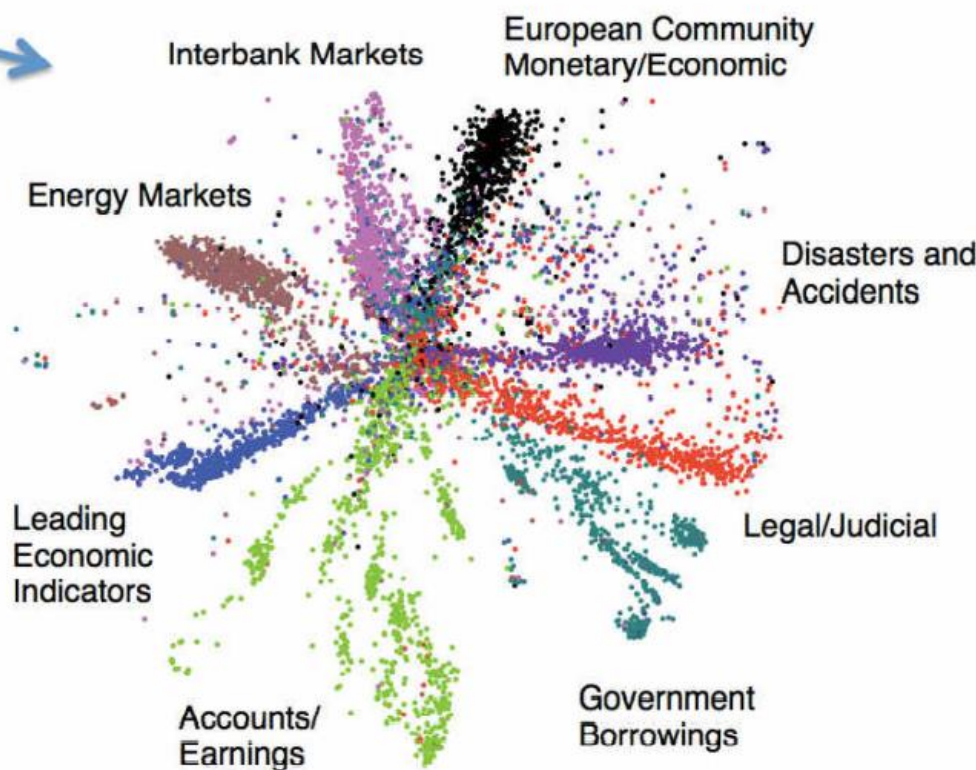
# Example application of DBN: Clustering of documents in database

Model  $P(\text{document})$



Bag of words

Reuters dataset: 804,414  
newswire stories: **unsupervised**

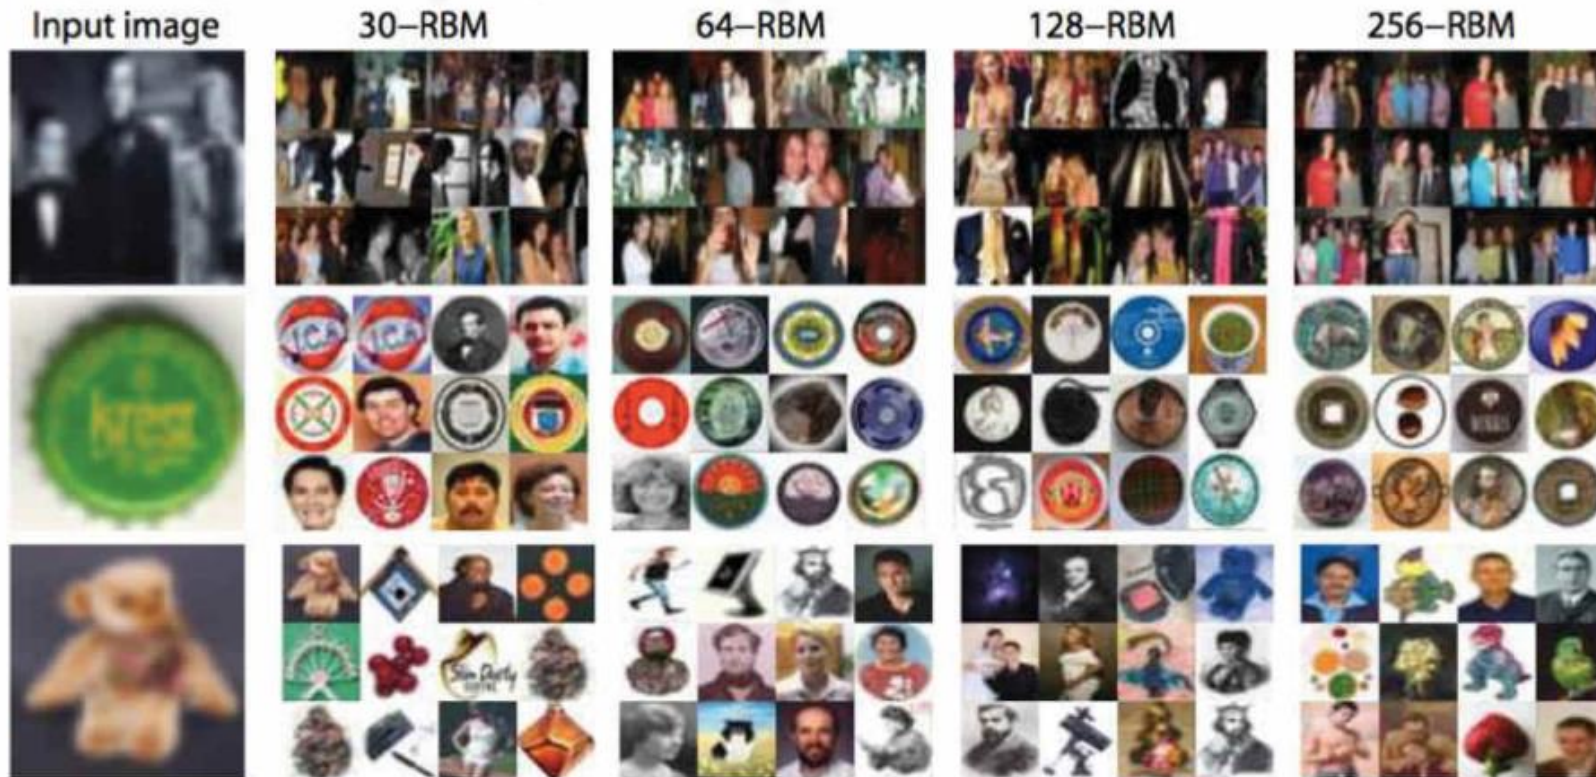


\*The goal of this model is to find the internal representation in all the types (classes) of documents provided by Reuters dataset



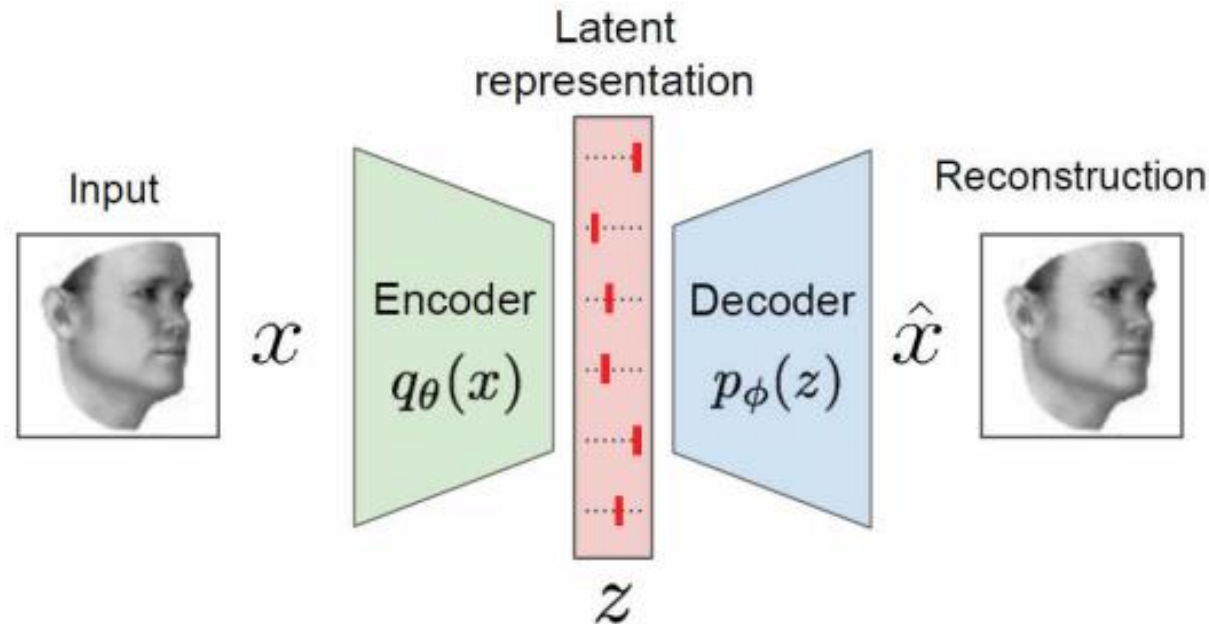
# Image Retrieval application example

- Map images in to binary codes fast retrieval



[Small codes, Torralba, Fergus, Weiss, CVPR 2008]

# Autoencoders

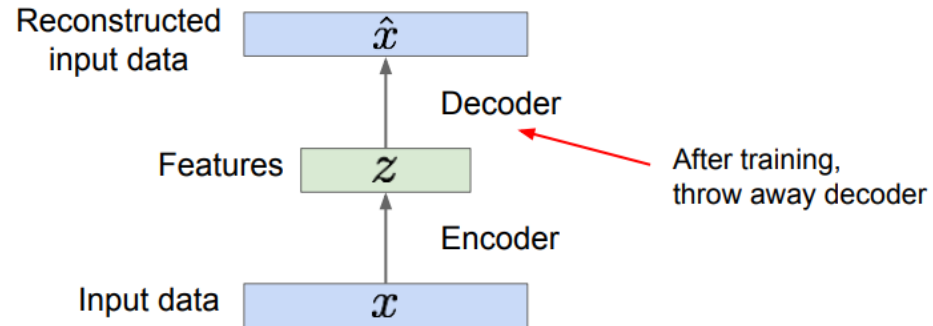


Learn  $q_{\theta}$  and  $p_{\phi}$  in order to minimize reconstruction cost:

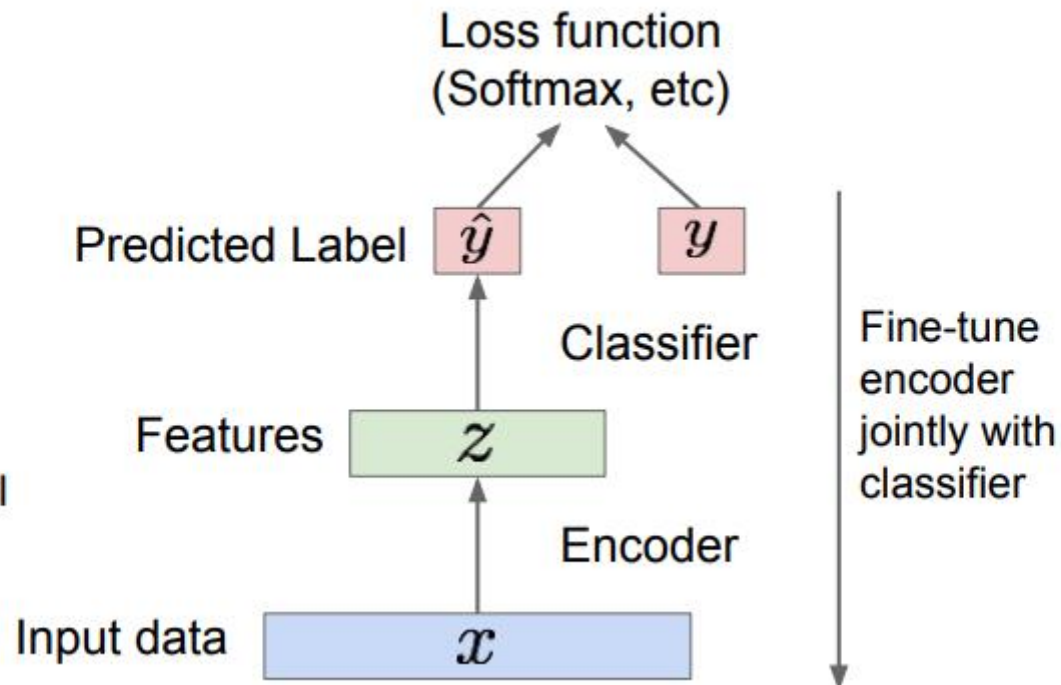
$$Q = \sum_k \|\hat{x}_k - x_k\|^2 = \sum_k \|p_{\phi}(q_{\theta}(x_k)) - x_k\|^2$$

→ unsupervised learning of latent variables,  
and of a generative model

# Use case of AutoEncoder



Encoder can be used to initialize a **supervised** model



bird plane  
dog deer truck

Train for final task  
(sometimes with small data)

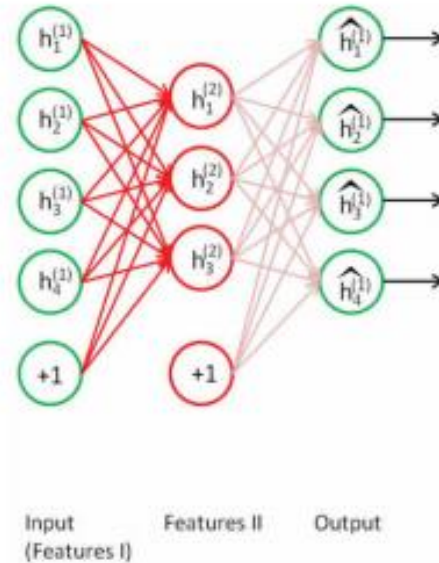
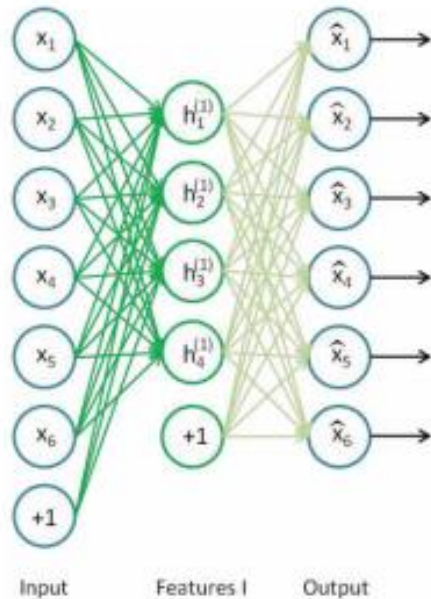




# Variants of autoencoders

- Denoising autoencoders
- Sparse autoencoders
- Stochastic autoencoders
- Contractive autoencoders
- VARIATIONAL autoencoders
- ...

# Training of stacked Autoencoders



etc...

## Greedy layerwise training:

for each layer  $k$ , use backpropagation to minimize

$$\| \mathbf{A}_k(\mathbf{h}^{(k)}) - \mathbf{h}^{(k)} \|^2 \quad (+ \text{ regularization cost } \lambda \sum_{ij} |W_{ij}|^2)$$

possibly + additional term for "sparsity"

# AutoEncoders (AE)

Unsupervised approach for learning a lower-dimensional feature representation from unlabeled training data

$z$  usually smaller than  $x$   
(dimensionality reduction)

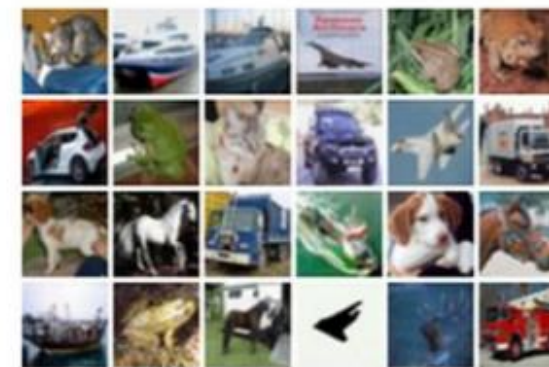
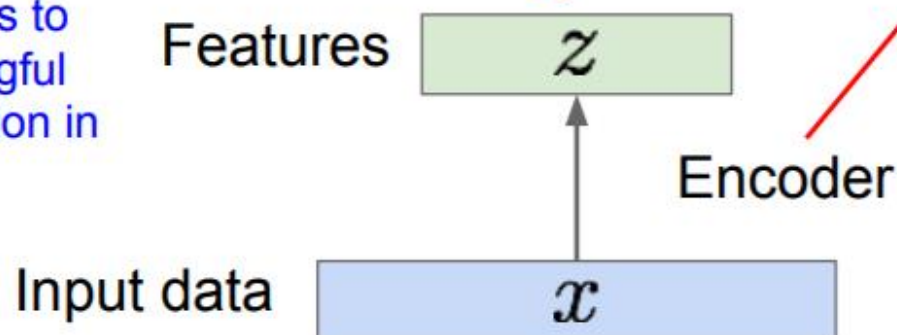
Q: Why dimensionality reduction?

A: Want features to capture meaningful factors of variation in data

**Originally:** Linear + nonlinearity (sigmoid)

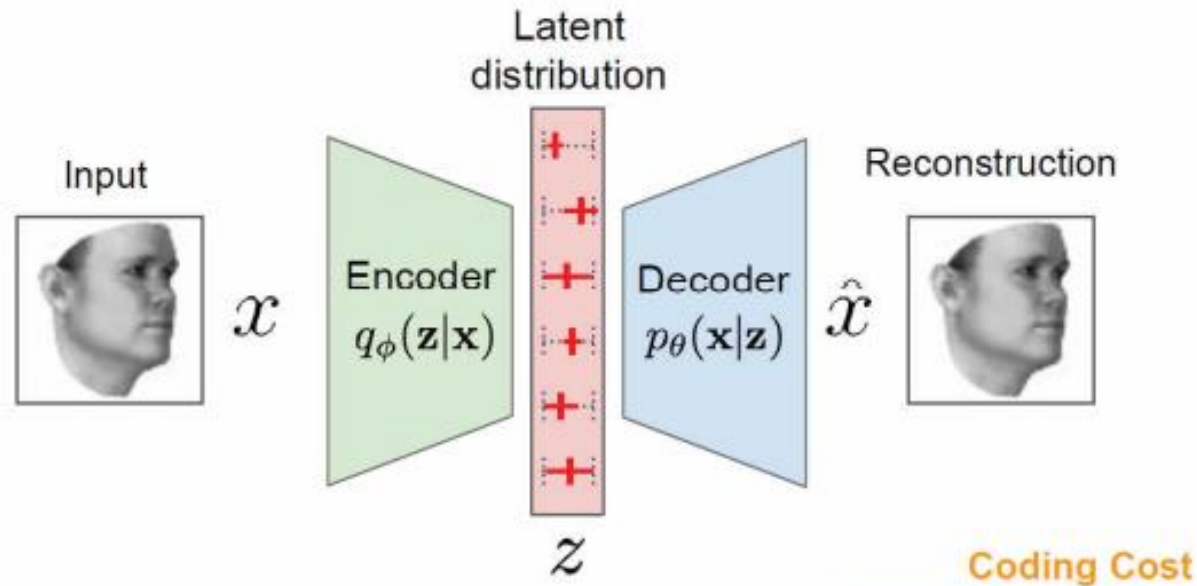
**Later:** Deep, fully-connected

**Later:** ReLU CNN



# Variational AutoEncoders (VAE)

Kingma et al. 2014  
Rezende et al. 2014



$$\mathcal{L}_{VAE}(\mathbf{x}; \theta, \phi) = \underbrace{\mathbb{E}_{q_{\phi}(\mathbf{z}|\mathbf{x})} [-\log p_{\theta}(\mathbf{x}|\mathbf{z})]}_{\text{Reconstruction cost}} + \underbrace{KL(q_{\phi}(\mathbf{z}|\mathbf{x}) || p(\mathbf{z}))}_{\text{Coding Cost}}$$

Slide: Irina Higgins, Loïc Matthey

**$KL$  = Kullback-Leibler divergence (a.k.a. ‘relative entropy’)**  
 **$KL(Q || P)$  measures how different are distributions**



Probabilistic spin on autoencoders - will let us sample from the model to generate data!

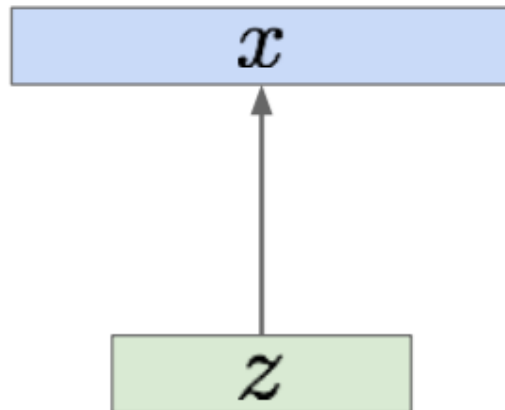
Assume training data  $\{x^{(i)}\}_{i=1}^N$  is generated from underlying unobserved (latent) representation  $\mathbf{z}$

Sample from  
true conditional

$$p_{\theta^*}(x \mid z^{(i)})$$

Sample from  
true prior

$$p_{\theta^*}(z)$$



**Intuition** (remember from autoencoders!):  
**x** is an image, **z** is latent factors used to  
generate **x**: attributes, orientation, etc.

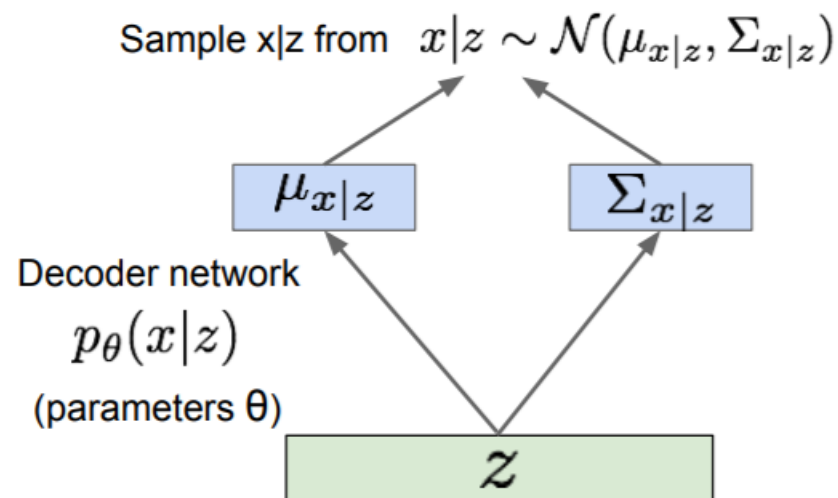
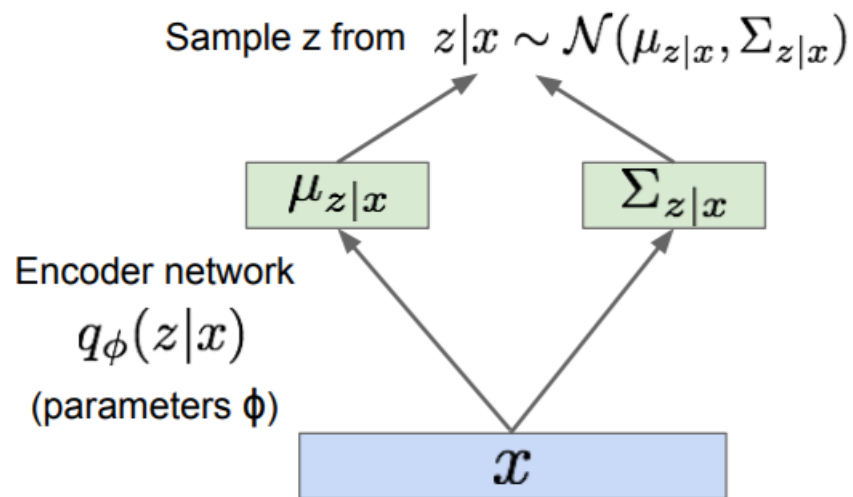
How should we represent this model?

Choose prior  $p(z)$  to be simple, e.g.  
Gaussian.

Conditional  $p(x|z)$  is complex (generates  
image) => represent with neural network

Data likelihood:  $p_{\theta}(x) = \int p_{\theta}(z) p_{\theta}(x|z) dz$

↑  
Intractable to compute  
 $p(x|z)$  for every  $z$ !





Now equipped with our encoder and decoder networks, let's work out the (log) data likelihood:

$$\begin{aligned}
 \log p_{\theta}(x^{(i)}) &= \mathbf{E}_{z \sim q_{\phi}(z|x^{(i)})} \left[ \log p_{\theta}(x^{(i)}) \right] && (p_{\theta}(x^{(i)})) \text{ Does not depend on } z \\
 &= \mathbf{E}_z \left[ \log \frac{p_{\theta}(x^{(i)} | z) p_{\theta}(z)}{p_{\theta}(z | x^{(i)})} \right] && (\text{Bayes' Rule}) \\
 &= \mathbf{E}_z \left[ \log \frac{p_{\theta}(x^{(i)} | z) p_{\theta}(z)}{p_{\theta}(z | x^{(i)})} \frac{q_{\phi}(z | x^{(i)})}{q_{\phi}(z | x^{(i)})} \right] && (\text{Multiply by constant}) \\
 &= \mathbf{E}_z \left[ \log p_{\theta}(x^{(i)} | z) \right] - \mathbf{E}_z \left[ \log \frac{q_{\phi}(z | x^{(i)})}{p_{\theta}(z)} \right] + \mathbf{E}_z \left[ \log \frac{q_{\phi}(z | x^{(i)})}{p_{\theta}(z | x^{(i)})} \right] && (\text{Logarithms}) \\
 &= \mathbf{E}_z \left[ \log p_{\theta}(x^{(i)} | z) \right] - D_{KL}(q_{\phi}(z | x^{(i)}) || p_{\theta}(z)) + D_{KL}(q_{\phi}(z | x^{(i)}) || p_{\theta}(z | x^{(i)}))
 \end{aligned}$$

↑

Decoder network gives  $p_{\theta}(x|z)$ , can compute estimate of this term through sampling. (Sampling differentiable through reparam. trick. see paper.)

↑

This KL term (between Gaussians for encoder and  $z$  prior) has nice closed-form solution!

↑

$p_{\theta}(z|x)$  intractable (saw earlier), can't compute this KL term :( But we know KL divergence always  $\geq 0$ .

# Generative Adversarial Network (GAN)

*[Introduced in 2014 by Ian Goodfellow et al.  
(incl. Yoshua Bengio) from University of Montreal]*

**Goal: generate « artificial » but credible examples**  
*credible = sampled from same probability distribution  $p(x)$*

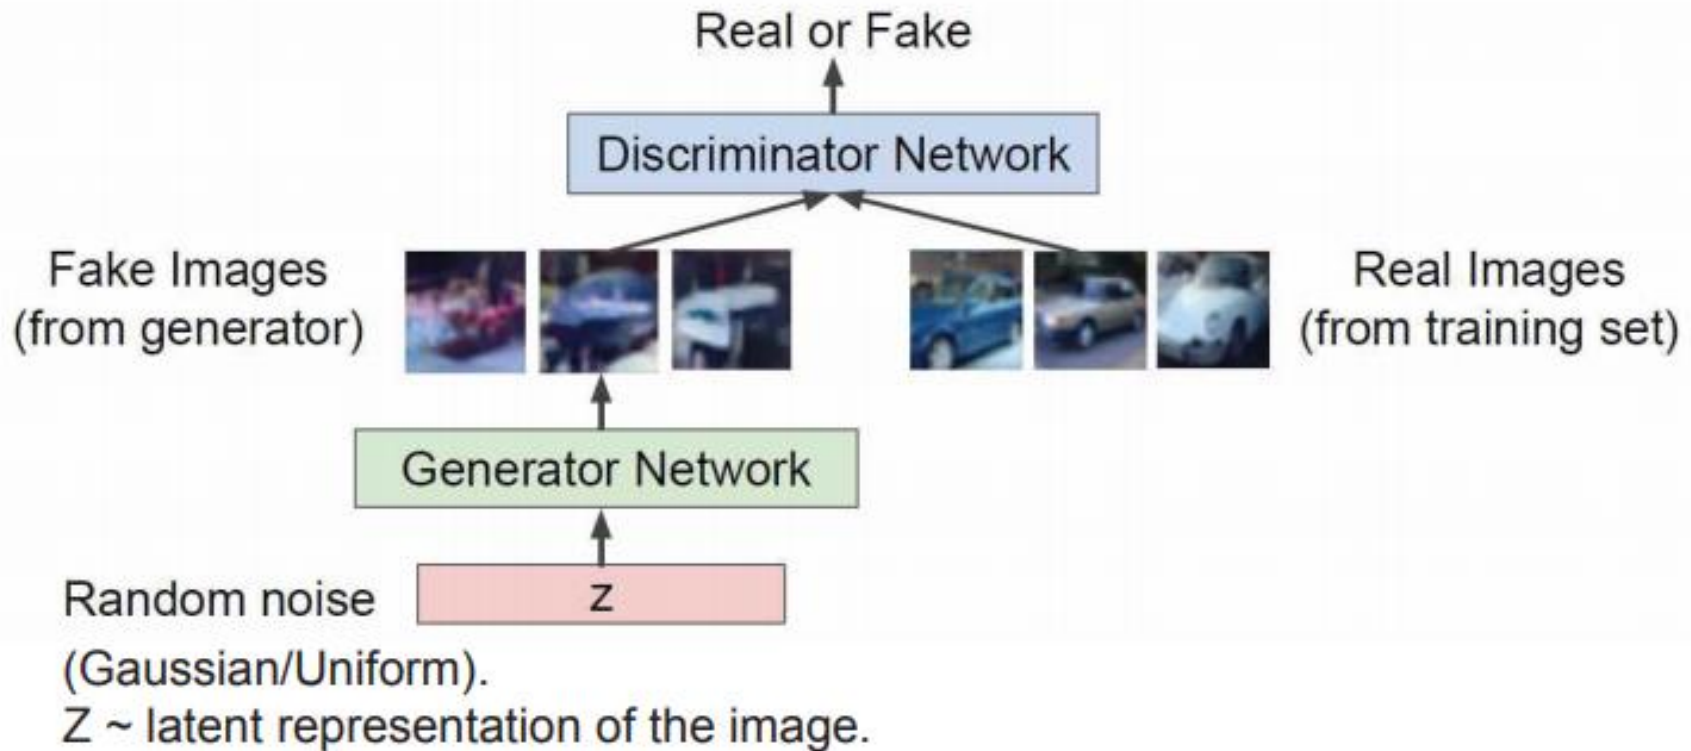
**Idea: instead of trying to explicitly estimate  $p(x)$ ,**

- 1. LEARN a transformation  $G$  from a simple and known distribution (e.g. random) into  $X$ ,**
- 2. then sampling  $z \rightarrow$  generate realistic samples  $G(z)$**

# GAN's architecture

**Generator network:** try to fool the discriminator by generating real-looking images

**Discriminator network:** try to distinguish between real and fake images





# GAN training: minimax two-player game!

$$\min_G \max_D V(D, G)$$

It is formulated as a **minimax game**, where:

- The Discriminator is trying to maximize its reward  $V(D, G)$
- The Generator is trying to minimize Discriminator's reward (or maximize its loss)

$$V(D, G) = \mathbb{E}_{x \sim p(x)} [\log D(x)] + \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

## Joint training of D and G

The Nash equilibrium of this particular game is achieved at:

- $P_{data}(x) = P_{gen}(x) \quad \forall x$
- $D(x) = \frac{1}{2} \quad \forall x$



# GAN training detail

**In practice, alternate Discriminator training (gradient *ascent*) and Generator training:**

**for** number of training iterations **do**

**for**  $k$  steps **do**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D_{\theta_d}(x^{(i)}) + \log(1 - D_{\theta_d}(G_{\theta_g}(z^{(i)}))) \right]$$

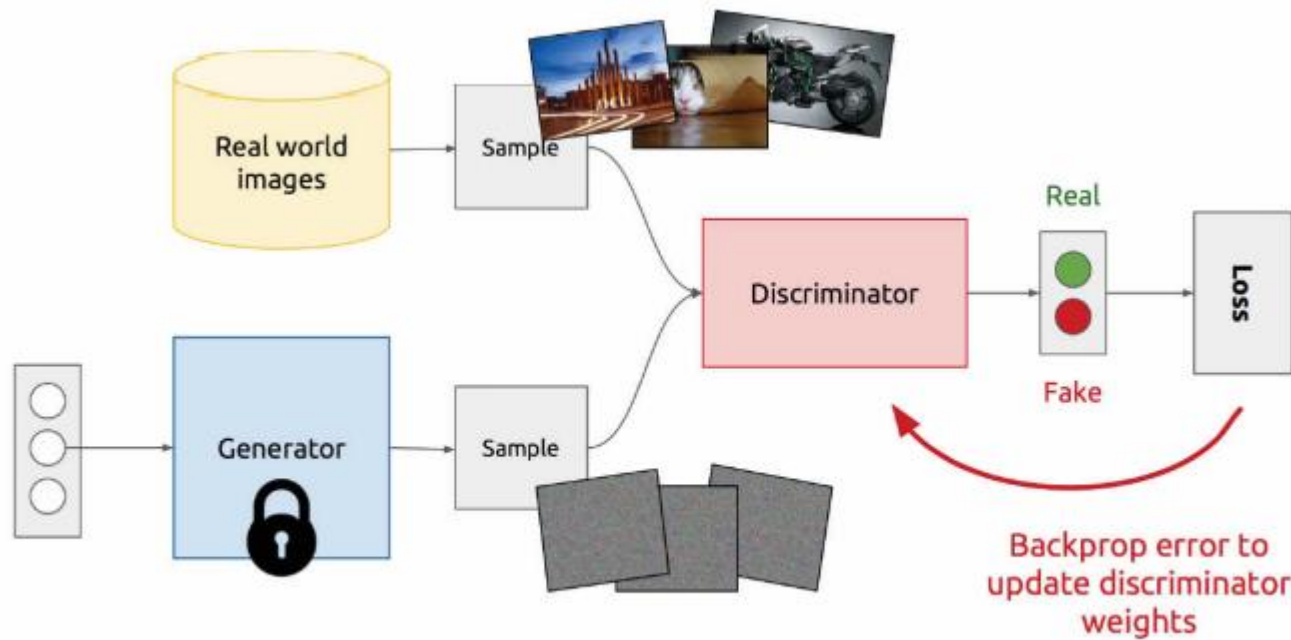
**end for**

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by ascending its stochastic gradient (improved objective):

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log(D_{\theta_d}(G_{\theta_g}(z^{(i)})))$$

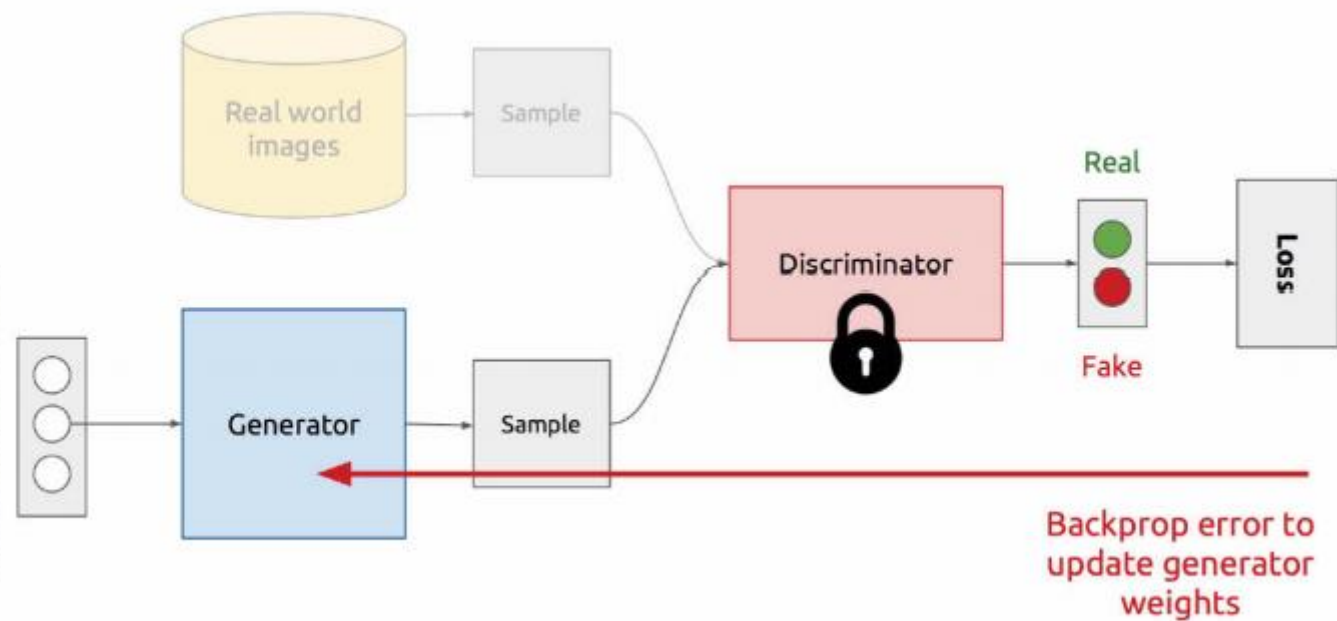
**end for**

Latent random variable

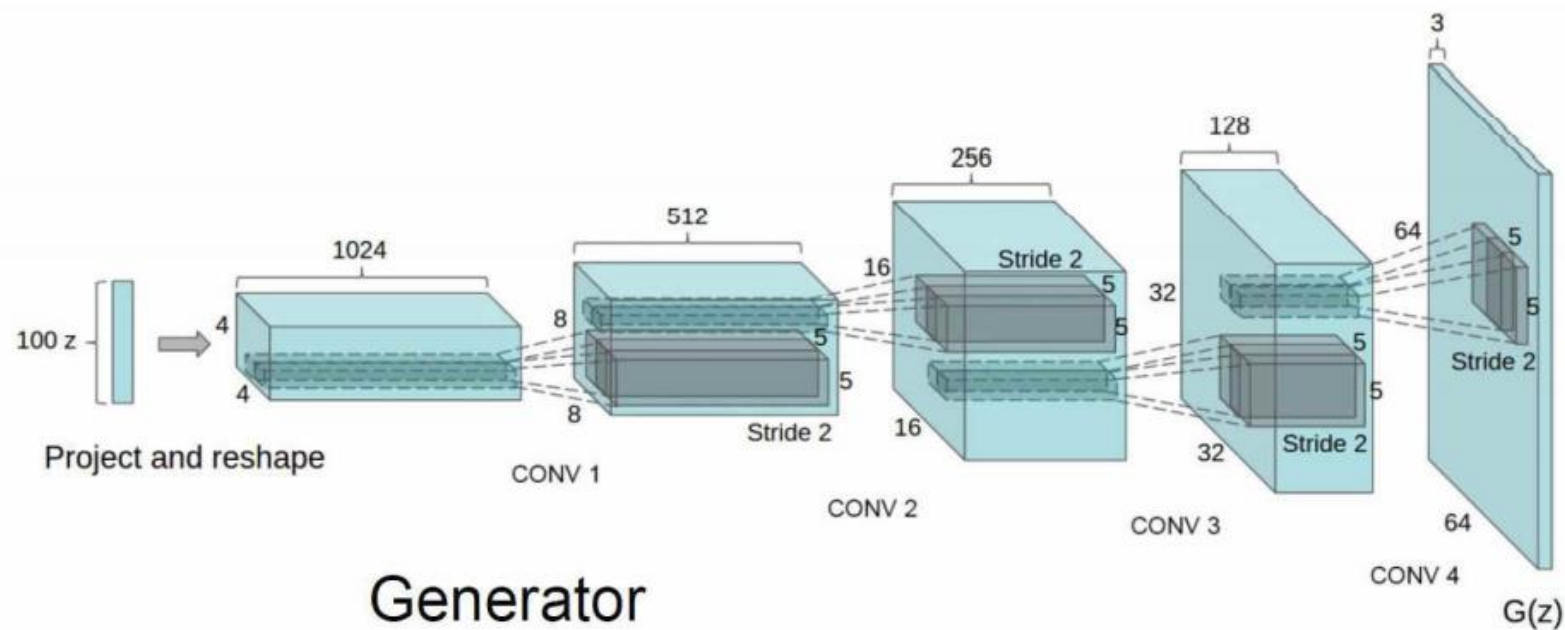


<https://www.slideshare.net/xavigiro/deep-learning-for-computer-vision-generative-models-and-adversarial-training-upc-2016>

Latent random variable



- Deep Convolutional Generative Adversarial Networks





# Fake images generated by DCGANs

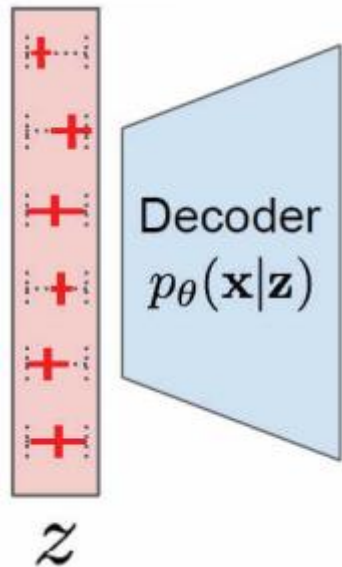
Samples  
from the  
model look  
amazing!

Radford et al,  
ICLR 2016



# Trajectory in latent space

- Walking in the latent space ( $z$ ) can help us understand the landscape of it as well as to reason if the model has learned relevant and interesting representations





# Sample code

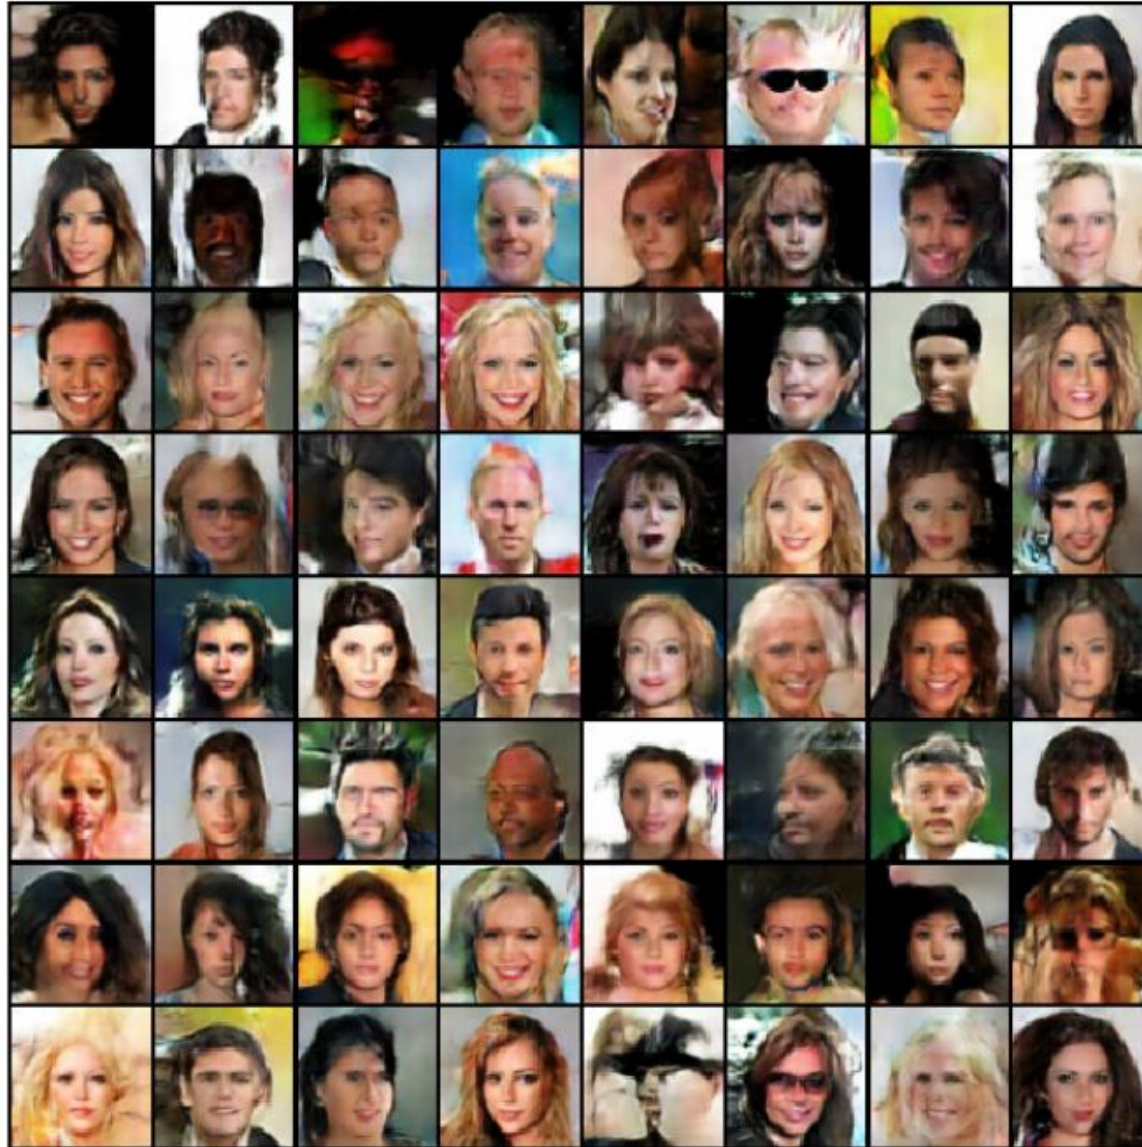
```
class Generator(nn.Module):
    def __init__(self, ngpu):
        super(Generator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is Z, going into a convolution
            nn.ConvTranspose2d( nz, ngf * 8, 4, 1, 0, bias=False),
            nn.BatchNorm2d(ngf * 8),
            nn.ReLU(True),
            # state size. (ngf*8) x 4 x 4
            nn.ConvTranspose2d(ngf * 8, ngf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 4),
            nn.ReLU(True),
            # state size. (ngf*4) x 8 x 8
            nn.ConvTranspose2d( ngf * 4, ngf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf * 2),
            nn.ReLU(True),
            # state size. (ngf*2) x 16 x 16
            nn.ConvTranspose2d( ngf * 2, ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),
            # state size. (ngf) x 32 x 32
            nn.ConvTranspose2d( ngf, nc, 4, 2, 1, bias=False),
            nn.Tanh()
            # state size. (nc) x 64 x 64
        )

    def forward(self, input):
        return self.main(input)
```

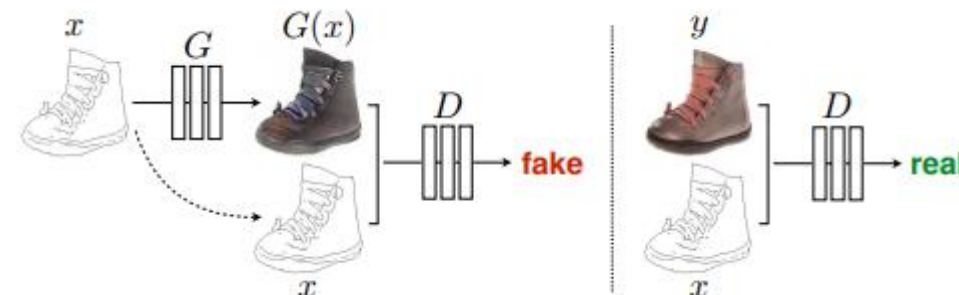
```
class Discriminator(nn.Module):
    def __init__(self, ngpu):
        super(Discriminator, self).__init__()
        self.ngpu = ngpu
        self.main = nn.Sequential(
            # input is (nc) x 64 x 64
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf) x 32 x 32
            nn.Conv2d(ndf, ndf * 2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 2),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*2) x 16 x 16
            nn.Conv2d(ndf * 2, ndf * 4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 4),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*4) x 8 x 8
            nn.Conv2d(ndf * 4, ndf * 8, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf * 8),
            nn.LeakyReLU(0.2, inplace=True),
            # state size. (ndf*8) x 4 x 4
            nn.Conv2d(ndf * 8, 1, 4, 1, 0, bias=False),
            nn.Sigmoid()
        )

    def forward(self, input):
        return self.main(input)
```

# Result of sample code (iteration 8000)



# Conditional GAN



- Training a conditional GAN to map edges  $\rightarrow$  photo. The discriminator,  $D$ , learns to classify between fake (synthesized by the generator) and real {edge, photo} tuples. The generator,  $G$ , learns to fool the discriminator. Unlike an unconditional GAN, both the generator and discriminator observe the input edge map

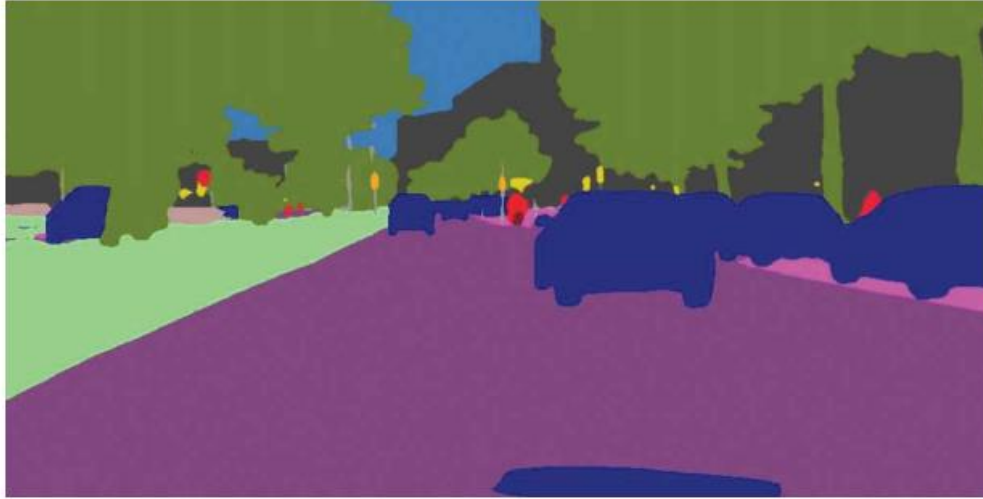


Interactive demo: <https://affinelayer.com/pixsrv/>

Isola, Phillip, et al. "Image-to-image translation with conditional adversarial networks." 2017.



# Video2video Synthesis

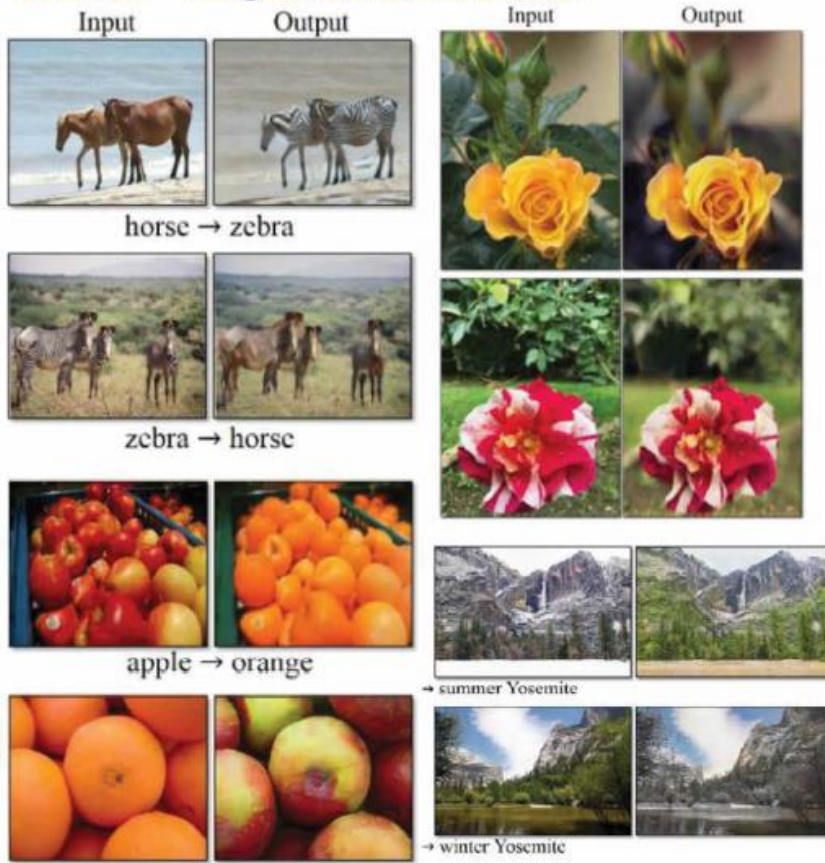


**"Video-to-Video Synthesis", NeurIPS'2018 [Nvidia+MIT]**  
*Using Generative Adversarial Network (GAN)*



# CycleGAN

## Source->Target domain transfer



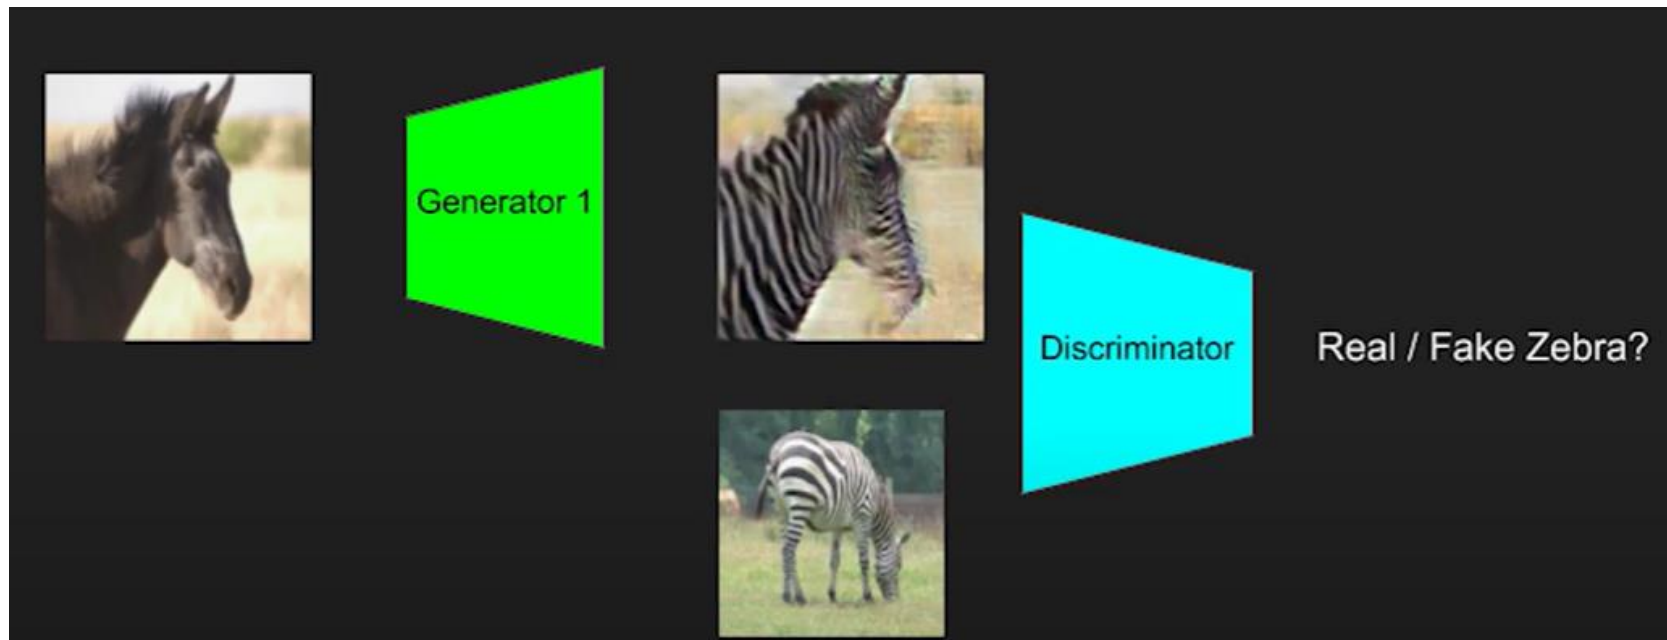
CycleGAN. Zhu et al. 2017.

- CycleGAN is proposed in 2017 (Jun-Yan Zhu, Taesung Park, et al. ) to deal with task with Unpaired Image-to-Image Translation
- The architecture use two generators and on discriminator.
- The Objectives are
  - Ensure the translated image looks like zebra
    - This is trained using the GAN objective with the discriminator
  - Ensure the translated image still looks mostly like the original
    - This is trained using a reconstruction objective with the second generator
    - This is the novel cycle-consistency loss





# CycleGAN



# “The GAN Zoo”

- GAN - Generative Adversarial Networks
- 3D-GAN - Learning a Probabilistic Latent Space of Object Shapes via 3D Generative-Adversarial Modeling
- acGAN - Face Aging With Conditional Generative Adversarial Networks
- AC-GAN - Conditional Image Synthesis With Auxiliary Classifier GANs
- AdaGAN - AdaGAN: Boosting Generative Models
- AEGAN - Learning Inverse Mapping by Autoencoder based Generative Adversarial Nets
- AffGAN - Amortised MAP Inference for Image Super-resolution
- AL-CGAN - Learning to Generate Images of Outdoor Scenes from Attributes and Semantic Layouts
- ALI - Adversarially Learned Inference
- AM-GAN - Generative Adversarial Nets with Labeled Data by Activation Maximization
- AnoGAN - Unsupervised Anomaly Detection with Generative Adversarial Networks to Guide Marker Discovery
- ArtGAN - ArtGAN: Artwork Synthesis with Conditional Categorical GANs
- b-GAN - b-GAN: Unified Framework of Generative Adversarial Networks
- Bayesian GAN - Deep and Hierarchical Implicit Models
- BEGAN - BEGAN: Boundary Equilibrium Generative Adversarial Networks
- BiGAN - Adversarial Feature Learning
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- CGAN - Conditional Generative Adversarial Nets
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks
- CCGAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CoGAN - Coupled Generative Adversarial Networks
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks
- DTN - Unsupervised Cross-Domain Image Generation
- DCGAN - Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
- DiscoGAN - Learning to Discover Cross-Domain Relations with Generative Adversarial Networks
- DR-GAN - Disentangled Representation Learning GAN for Pose-Invariant Face Recognition
- DualGAN - DualGAN: Unsupervised Dual Learning for Image-to-Image Translation
- EBGAN - Energy-based Generative Adversarial Network
- f-GAN - f-GAN: Training Generative Neural Samplers using Variational Divergence Minimization
- FF-GAN - Towards Large-Pose Face Frontalization in the Wild
- GAWWN - Learning What and Where to Draw
- GeneGAN - GeneGAN: Learning Object Transfiguration and Attribute Subspace from Unpaired Data
- Geometric GAN - Geometric GAN
- GoGAN - Gang of GANs: Generative Adversarial Networks with Maximum Margin Ranking
- GP-GAN - GP-GAN: Towards Realistic High-Resolution Image Blending
- IAN - Neural Photo Editing with Introspective Adversarial Networks
- iGAN - Generative Visual Manipulation on the Natural Image Manifold
- IcGAN - Invertible Conditional GANs for image editing
- ID-CGAN - Image De-raining Using a Conditional Generative Adversarial Network
- Improved GAN - Improved Techniques for Training GANs
- InfoGAN - InfoGAN: Interpretable Representation Learning by Information Maximizing Generative Adversarial Nets
- LAGAN - Learning Particle Physics by Example: Location-Aware Generative Adversarial Networks for Physics Synthesis
- LAPGAN - Deep Generative Image Models using a Laplacian Pyramid of Adversarial Networks

<https://github.com/hindupuravinash/the-gan-zoo>

# GANs short summary

- Don't work with an explicit density function
- Take game-theoretic approach: learn to generate from training distribution through 2-player game
- Pros:
  - Beautiful, state-of-the-art samples!
- Cons:
  - Trickier / more unstable to train
  - Can't solve inference queries such as  $p(x)$ ,  $p(z|x)$
- Active areas of research: -
  - Better loss functions, more stable training (Wasserstein GAN, LSGAN, many others)
  - Conditional GANs, GANs for all kinds of applications

# Summary

- Intrinsically UNSUPERVISED
  - can be used on UNLABELLED DATA
- Impressive results in Image Retrieval
- DBN/DBM/VAE = Generative probabilistic models
- GAN = most promising generative model, with already many remarkable & exciting applications
- Strong potential for enhancement of datasets and for ultra-realistic synthetic data
- Interest for "creative« /artistic computing?

# Any QUESTIONS ?